



AWS Solutions Architect Associate

Session 601

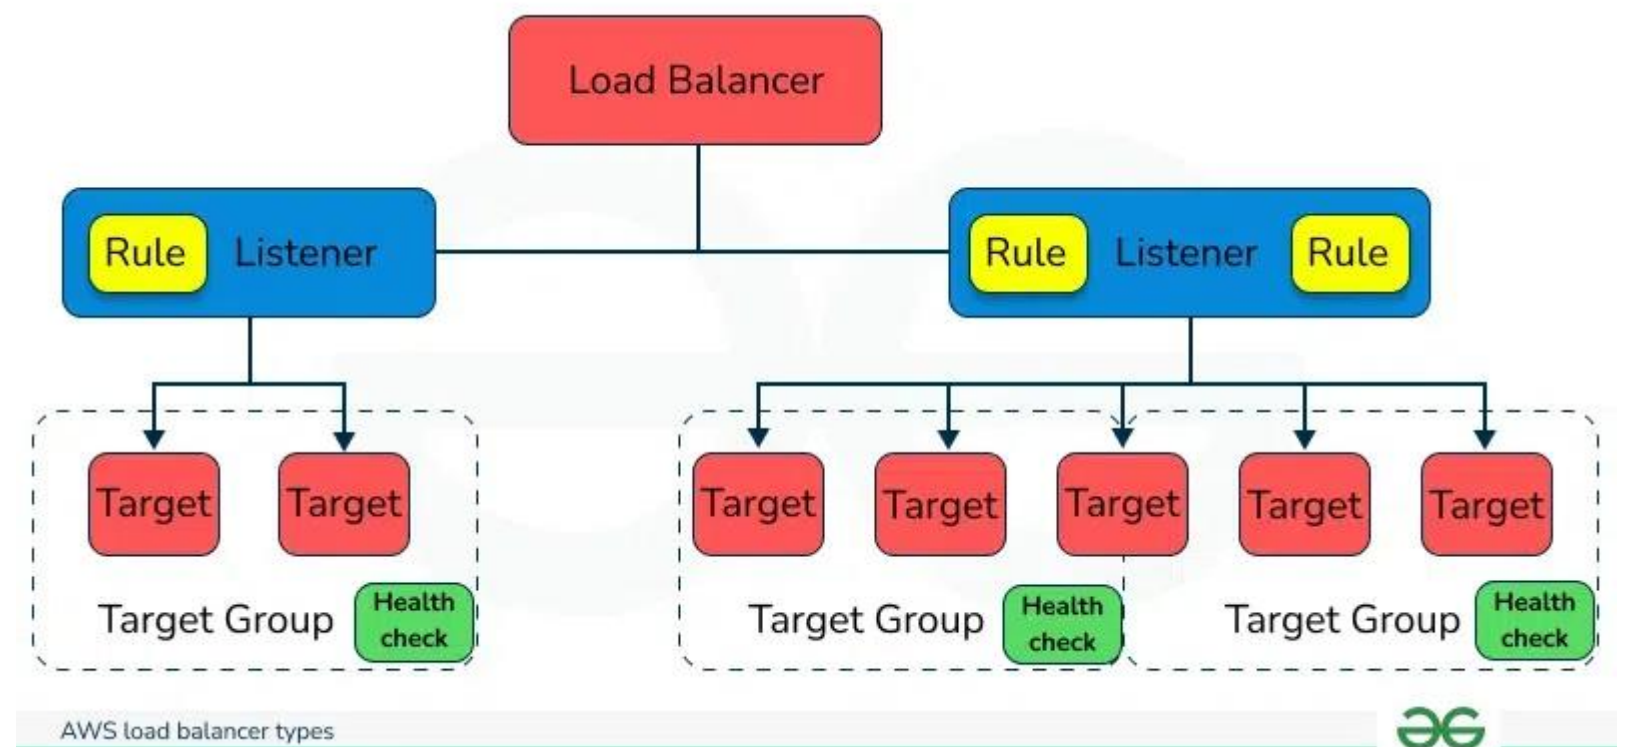
**Compute: ELB and Mgmt
&Gov: Cloudwatch**

July/2024



Elastic Load Balancing

A **managed load balancing service** that distributes incoming application traffic across multiple Amazon EC2 instances, containers, IP addresses and Lambda Functions.






Elastic Load
Balancing

- Uses HTTP, HTTPS, TCP and SSL (secure TCP) protocols.
- Can be external or internal facing
- Each load balancer is given a DNS name
- Recognizes and responds to healthy instances



High
availability




Health
checks

Protocol,
port, path



Security
features




TLS
termination



If you need to **remove an instance** from your production fleet, but **don't want to affect your users**:


Affected backend instances will complete requests in progress before deregistration



Enable connection draining



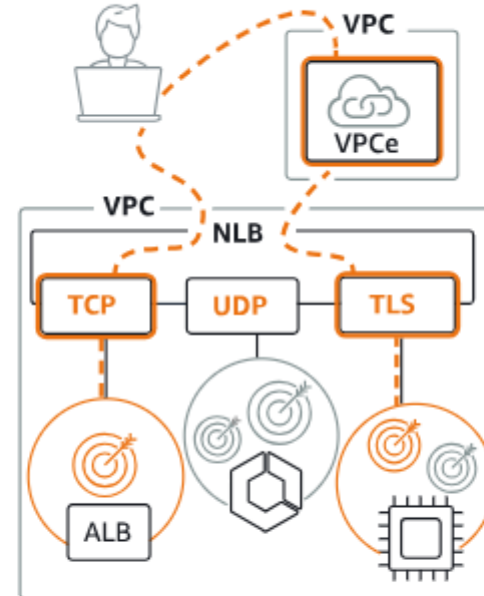
Application Load Balancer [Info](#)



Choose an Application Load Balancer when you need a flexible feature set for your applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing and visibility features targeted at application architectures, including microservices and containers.

Create

Network Load Balancer [Info](#)



Choose a Network Load Balancer when you need ultra-high performance, TLS offloading at scale, centralized certificate deployment, support for UDP, and static IP addresses for your applications. Operating at the connection level, Network Load Balancers are capable of handling millions of requests per second securely while maintaining ultra-low latencies.

Gateway Load Balancer [Info](#)




Choose a Gateway Load Balancer when you need to deploy and manage a fleet of third-party virtual appliances that support GENEVE. These appliances enable you to improve security, compliance, and policy controls.

Create

▼ Classic Load Balancer - *previous generation*

Classic Load Balancer [Info](#)

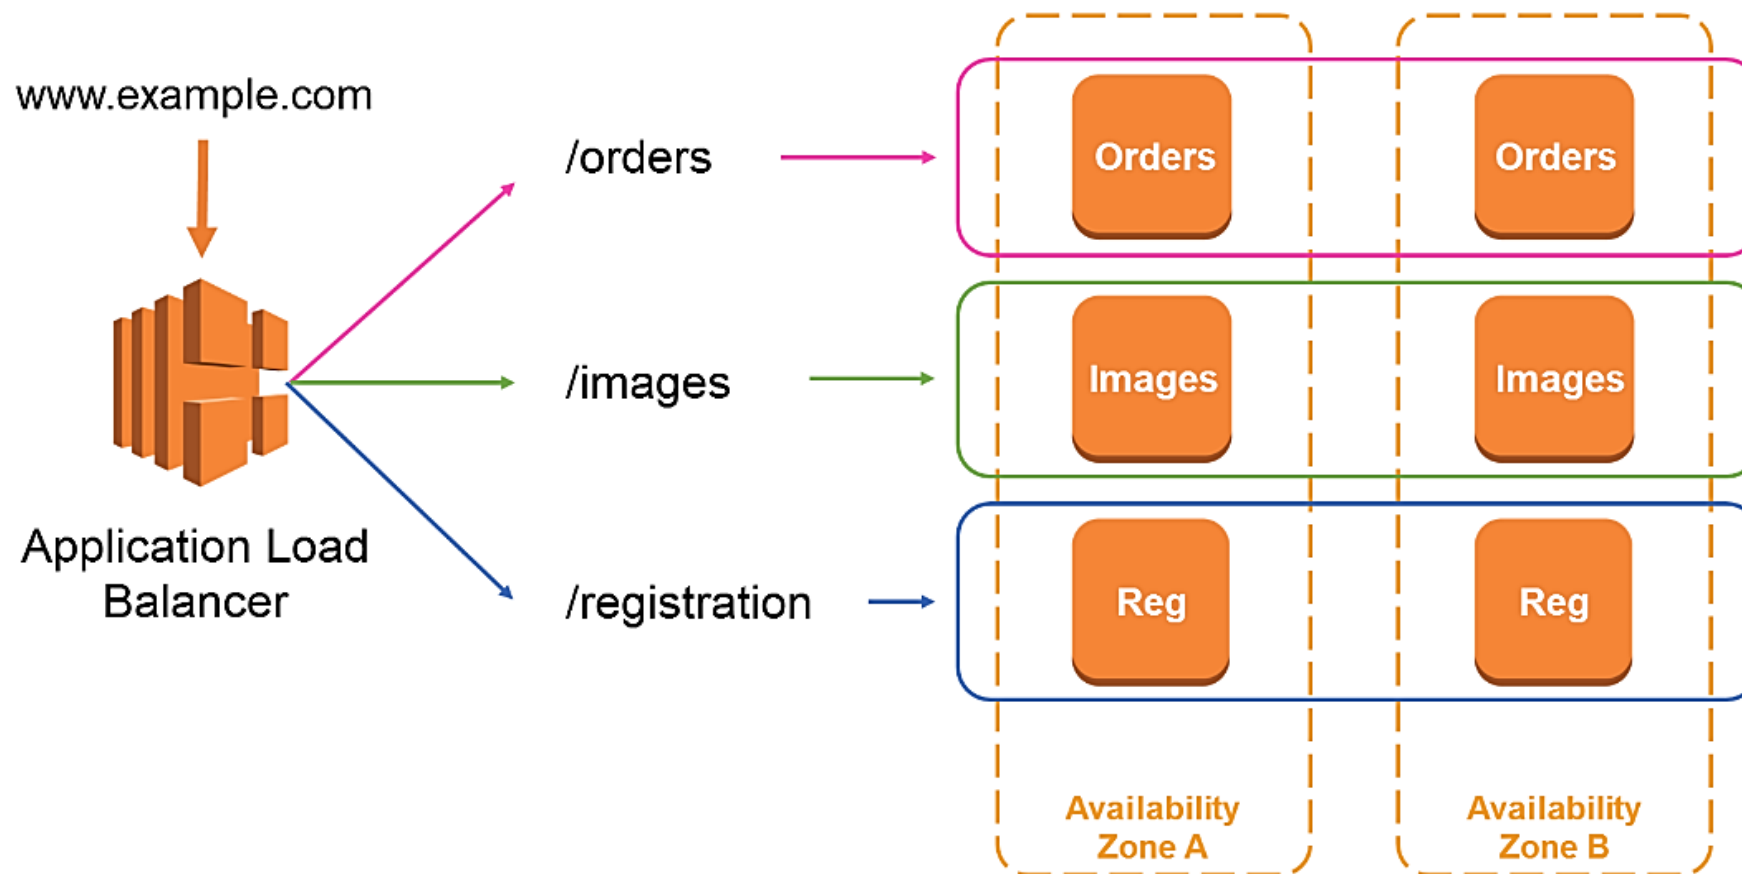


Choose a Classic Load Balancer when you have an existing application running in the EC2-Classic network.

Create



Example of ALB Feature: Path Routing

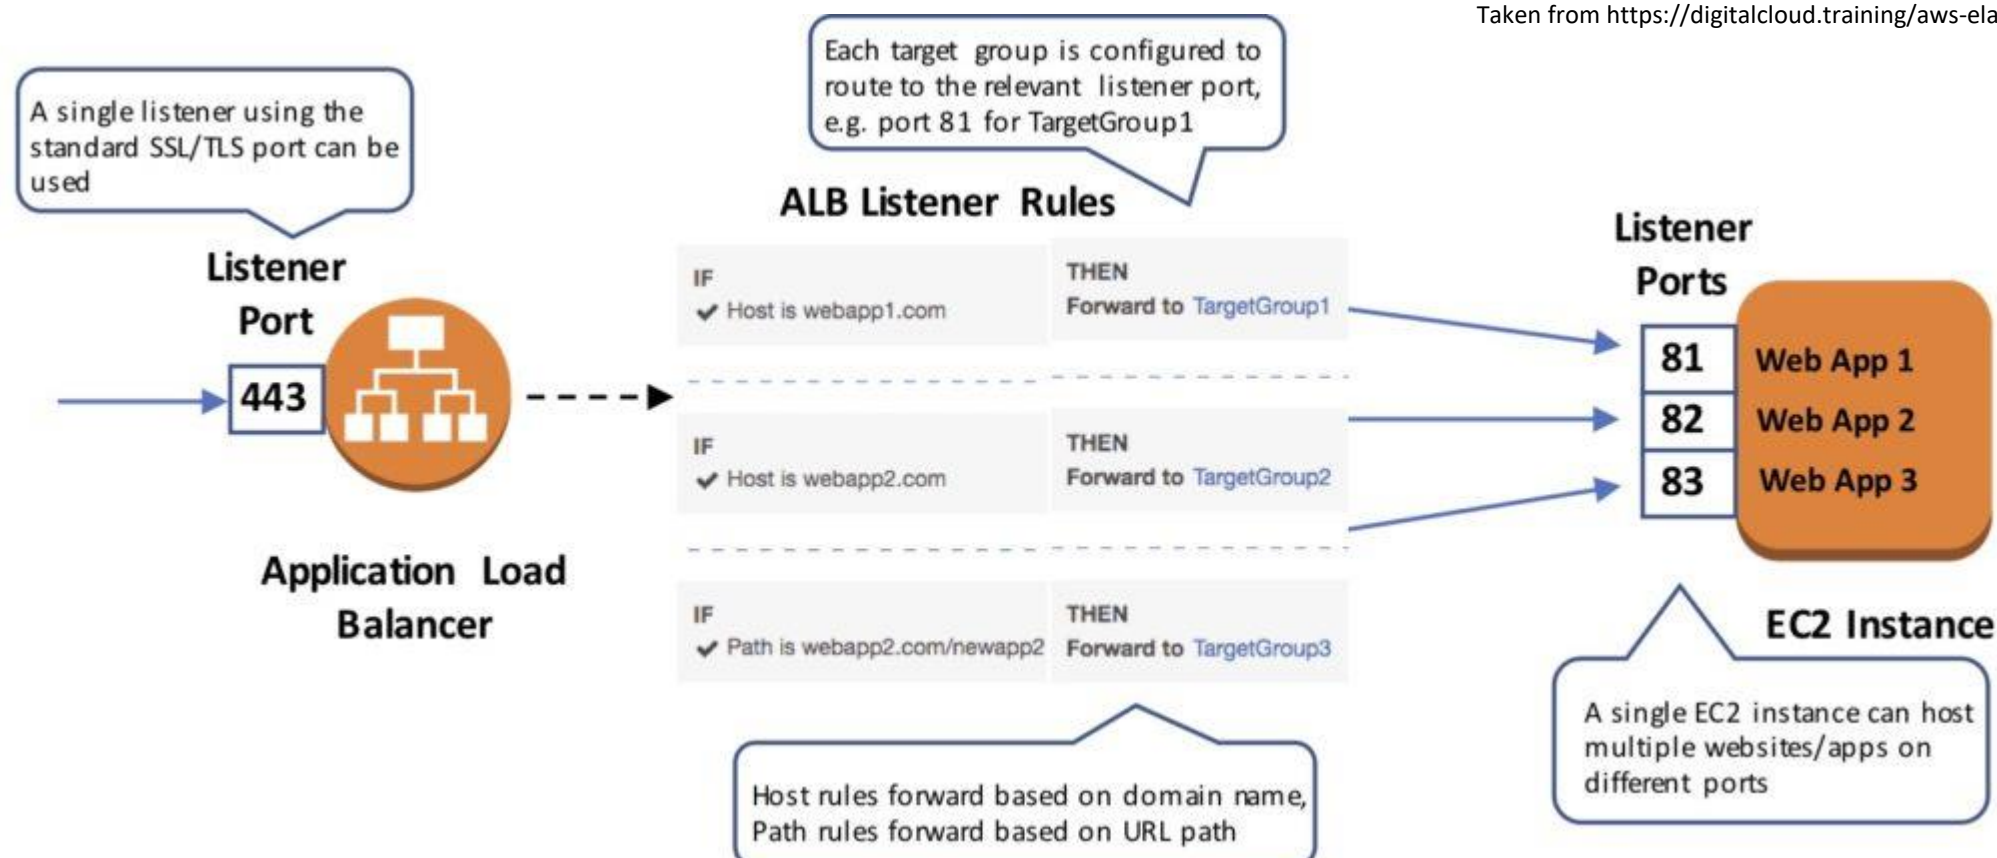



I: Target Groups: Type (Instance/IP), Port, Target Registration

II: Load Balancer: Type and Listener.

Note: Can be done using Infrastructure or Kubernetes Commands.


Taken from <https://digitalcloud.training/aws-elastic-load-balancing-aws-elb/> (18/07/2024)





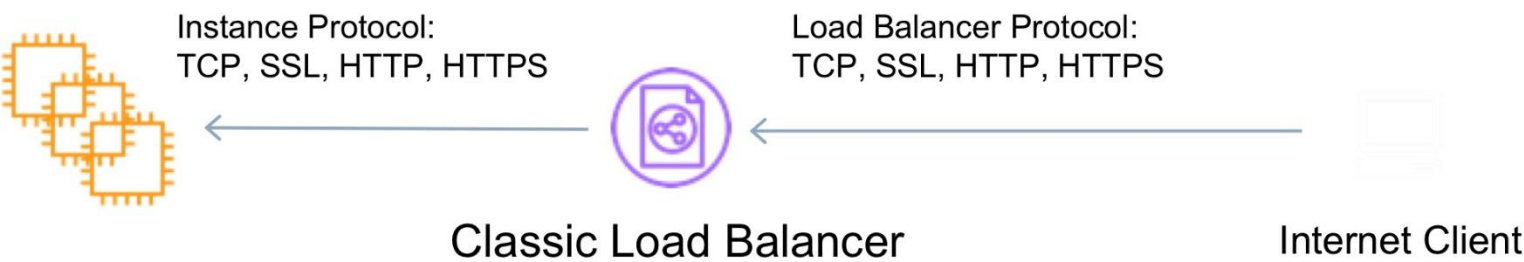
Application Load Balancer

- Operates at the request level
- Routes based on the content of the request (layer 7)
- Supports path-based routing, host-based routing, query string parameter-based routing, and source IP address-based routing
- Supports IP addresses, Lambda Functions and containers as targets



Network Load Balancer

- Operates at the connection level
- Routes connections based on IP protocol data (layer 4)
- Offers ultra high performance, low latency and TLS offloading at scale
- Can have static IP / Elastic IP
- Supports UDP and static IP addresses as targets



Classic Load Balancer

- Old generation; not recommended for new applications
- Performs routing at Layer 4 and Layer 7
- Use for existing applications running in EC2-Classic




Your application can **recover from a failure or roll over to a secondary source** within an **acceptable** amount of degraded performance time.

Percent of Uptime	Max Downtime per Year	Equivalent Downtime per Day
90%	36.5 days	2.4 hrs
99%	3.65 days	14 min
99.9%	8.76 hrs	86 sec
99.99%	52.6 min	8.6 sec
99.999%	5.25 min	.86 sec



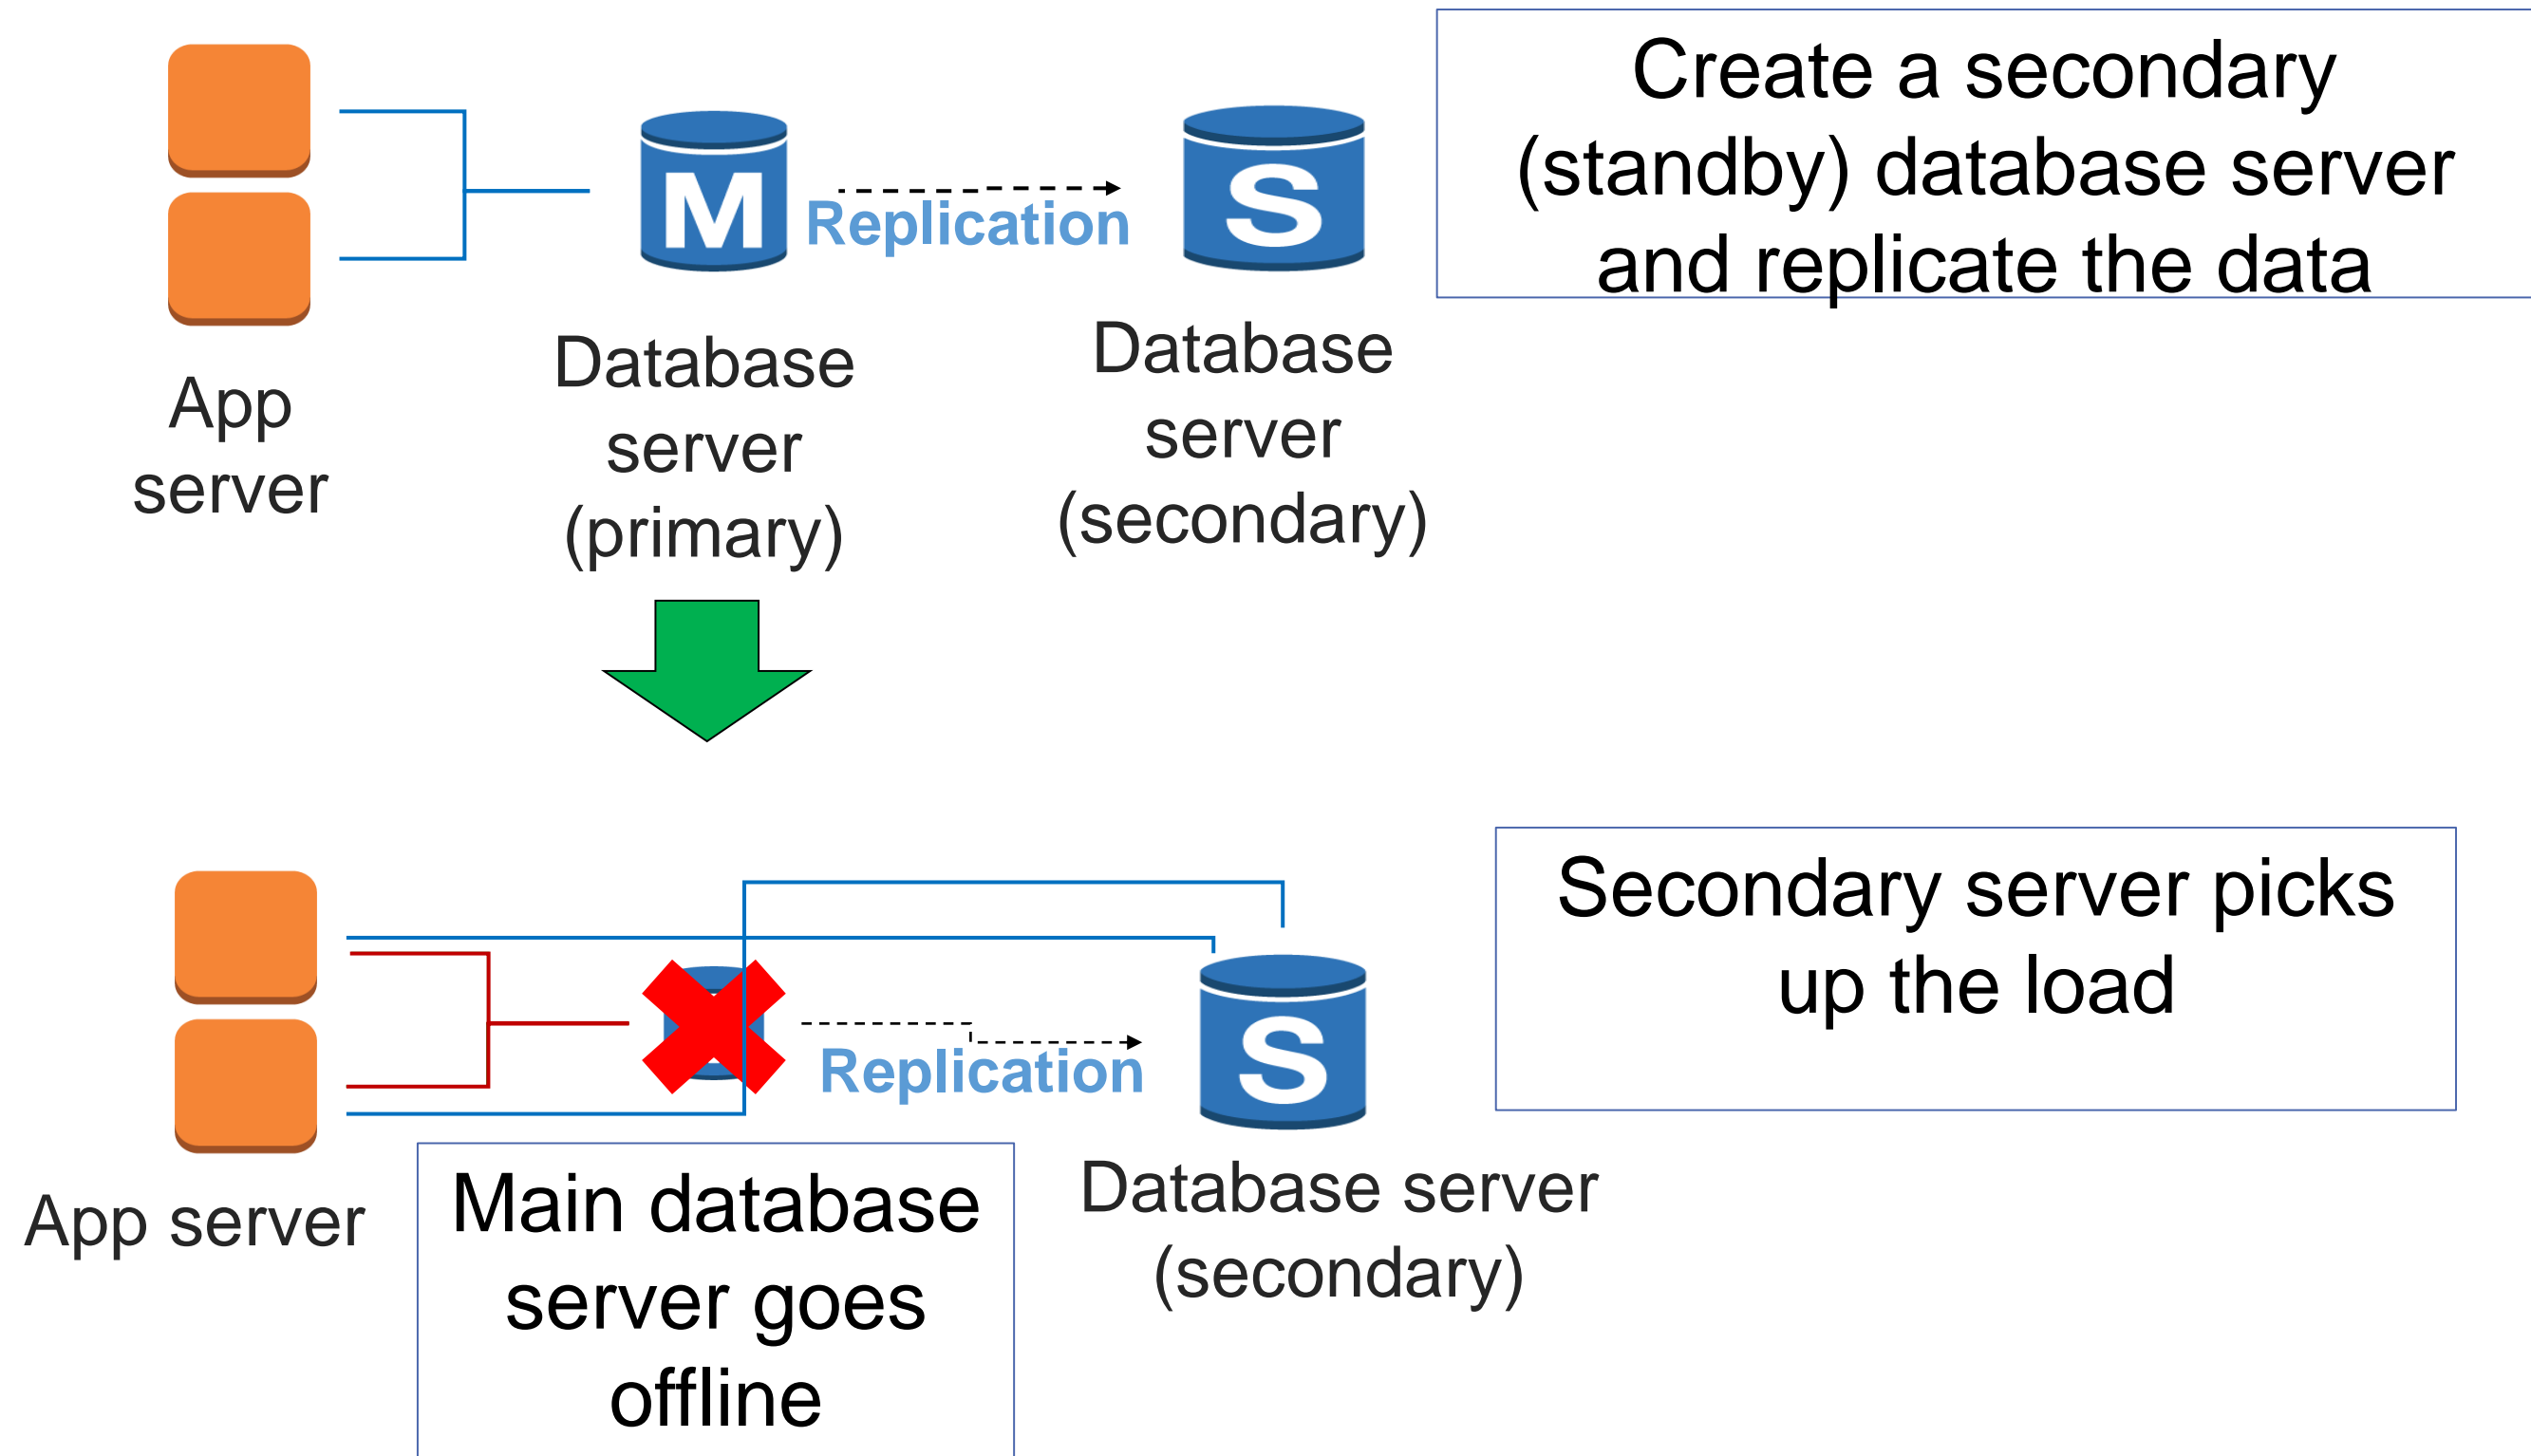
*Assume everything
fails, and design
backward*

Implement redundancy where possible
in order to prevent single failures from
bringing down an entire system.






Best practice

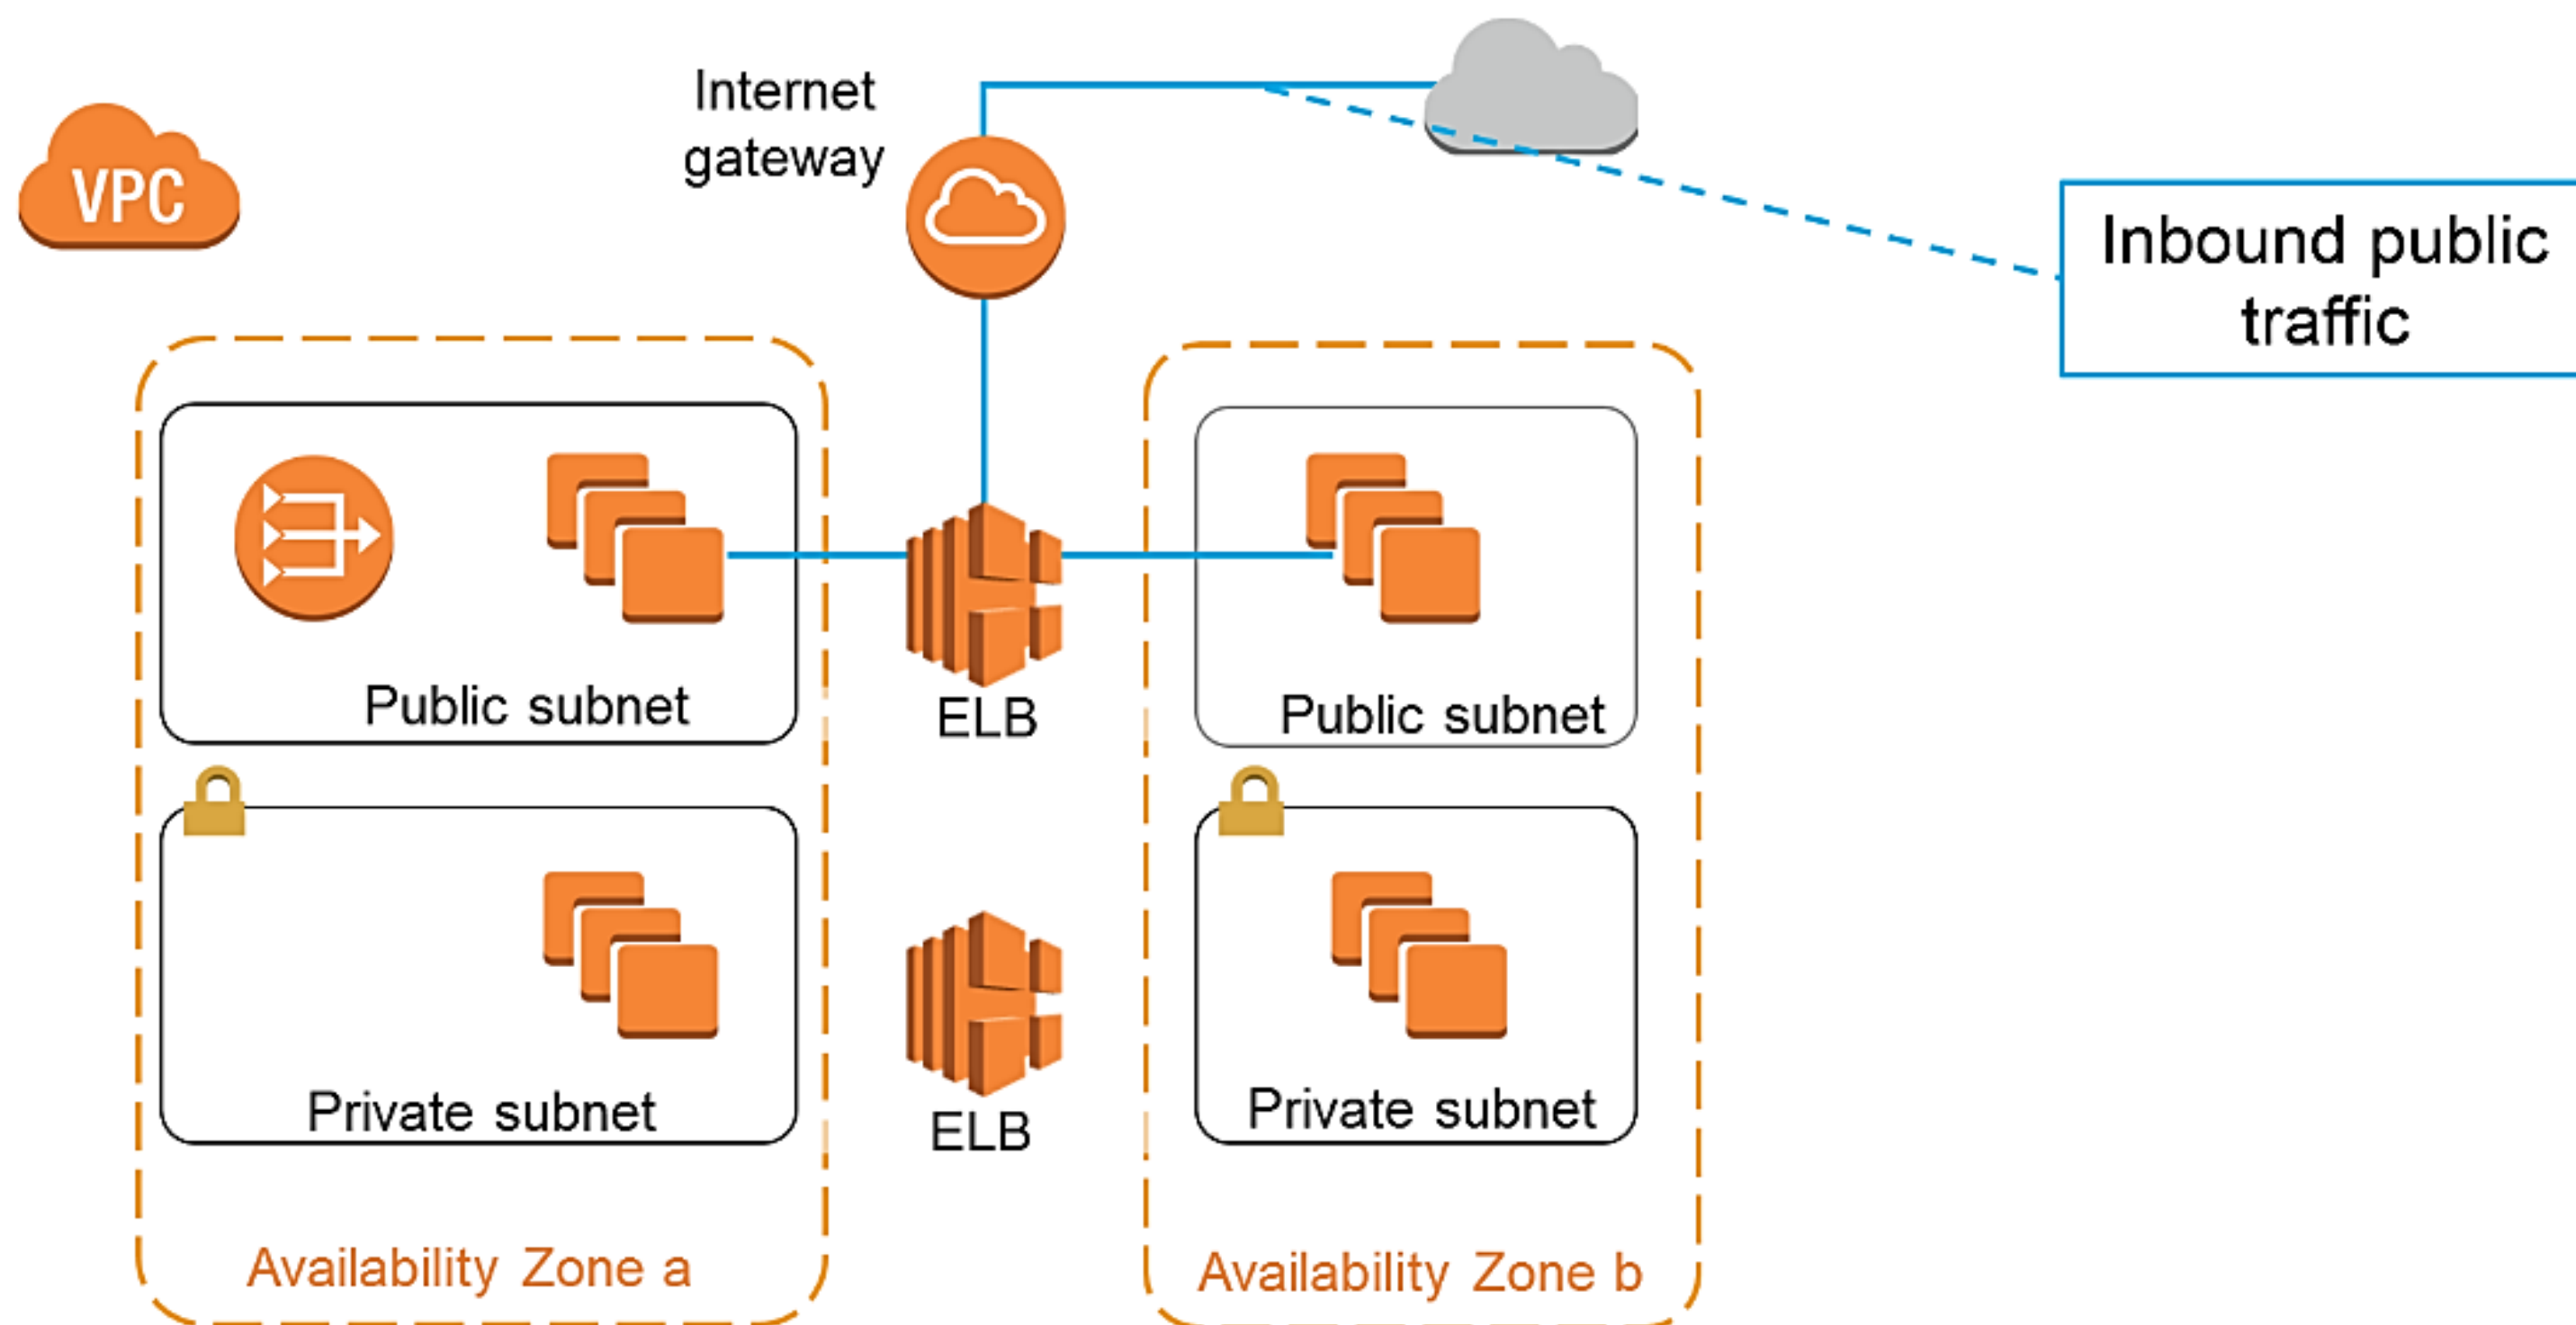


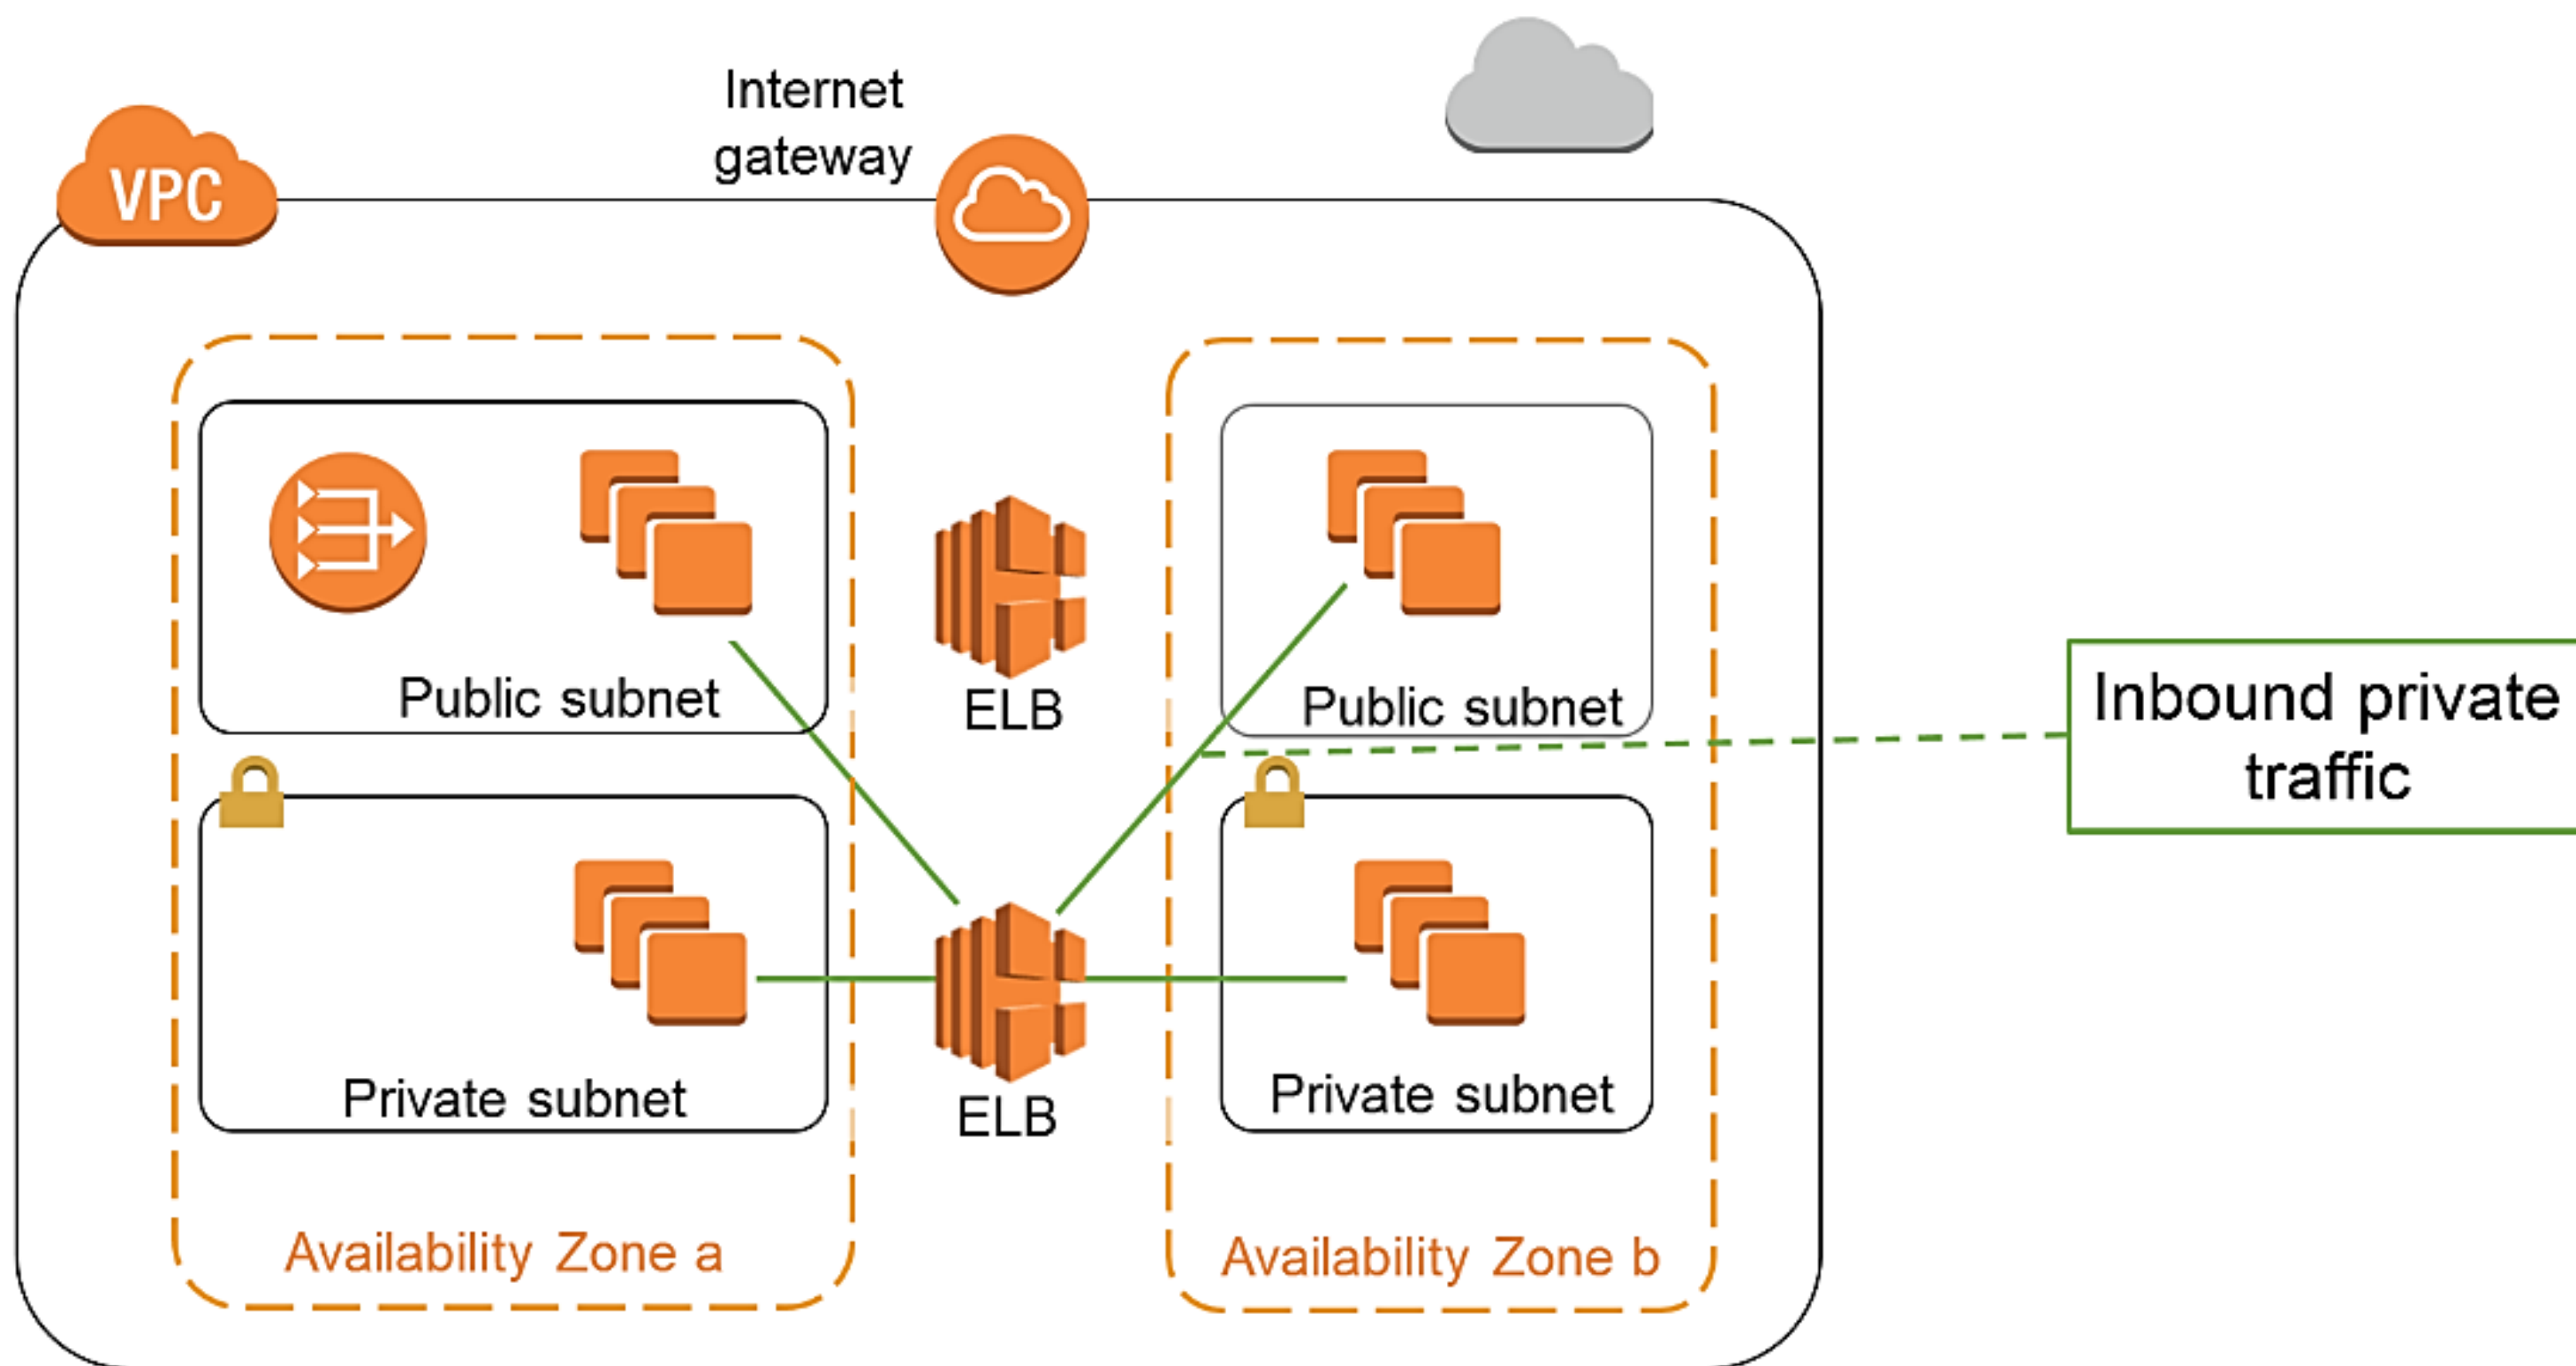


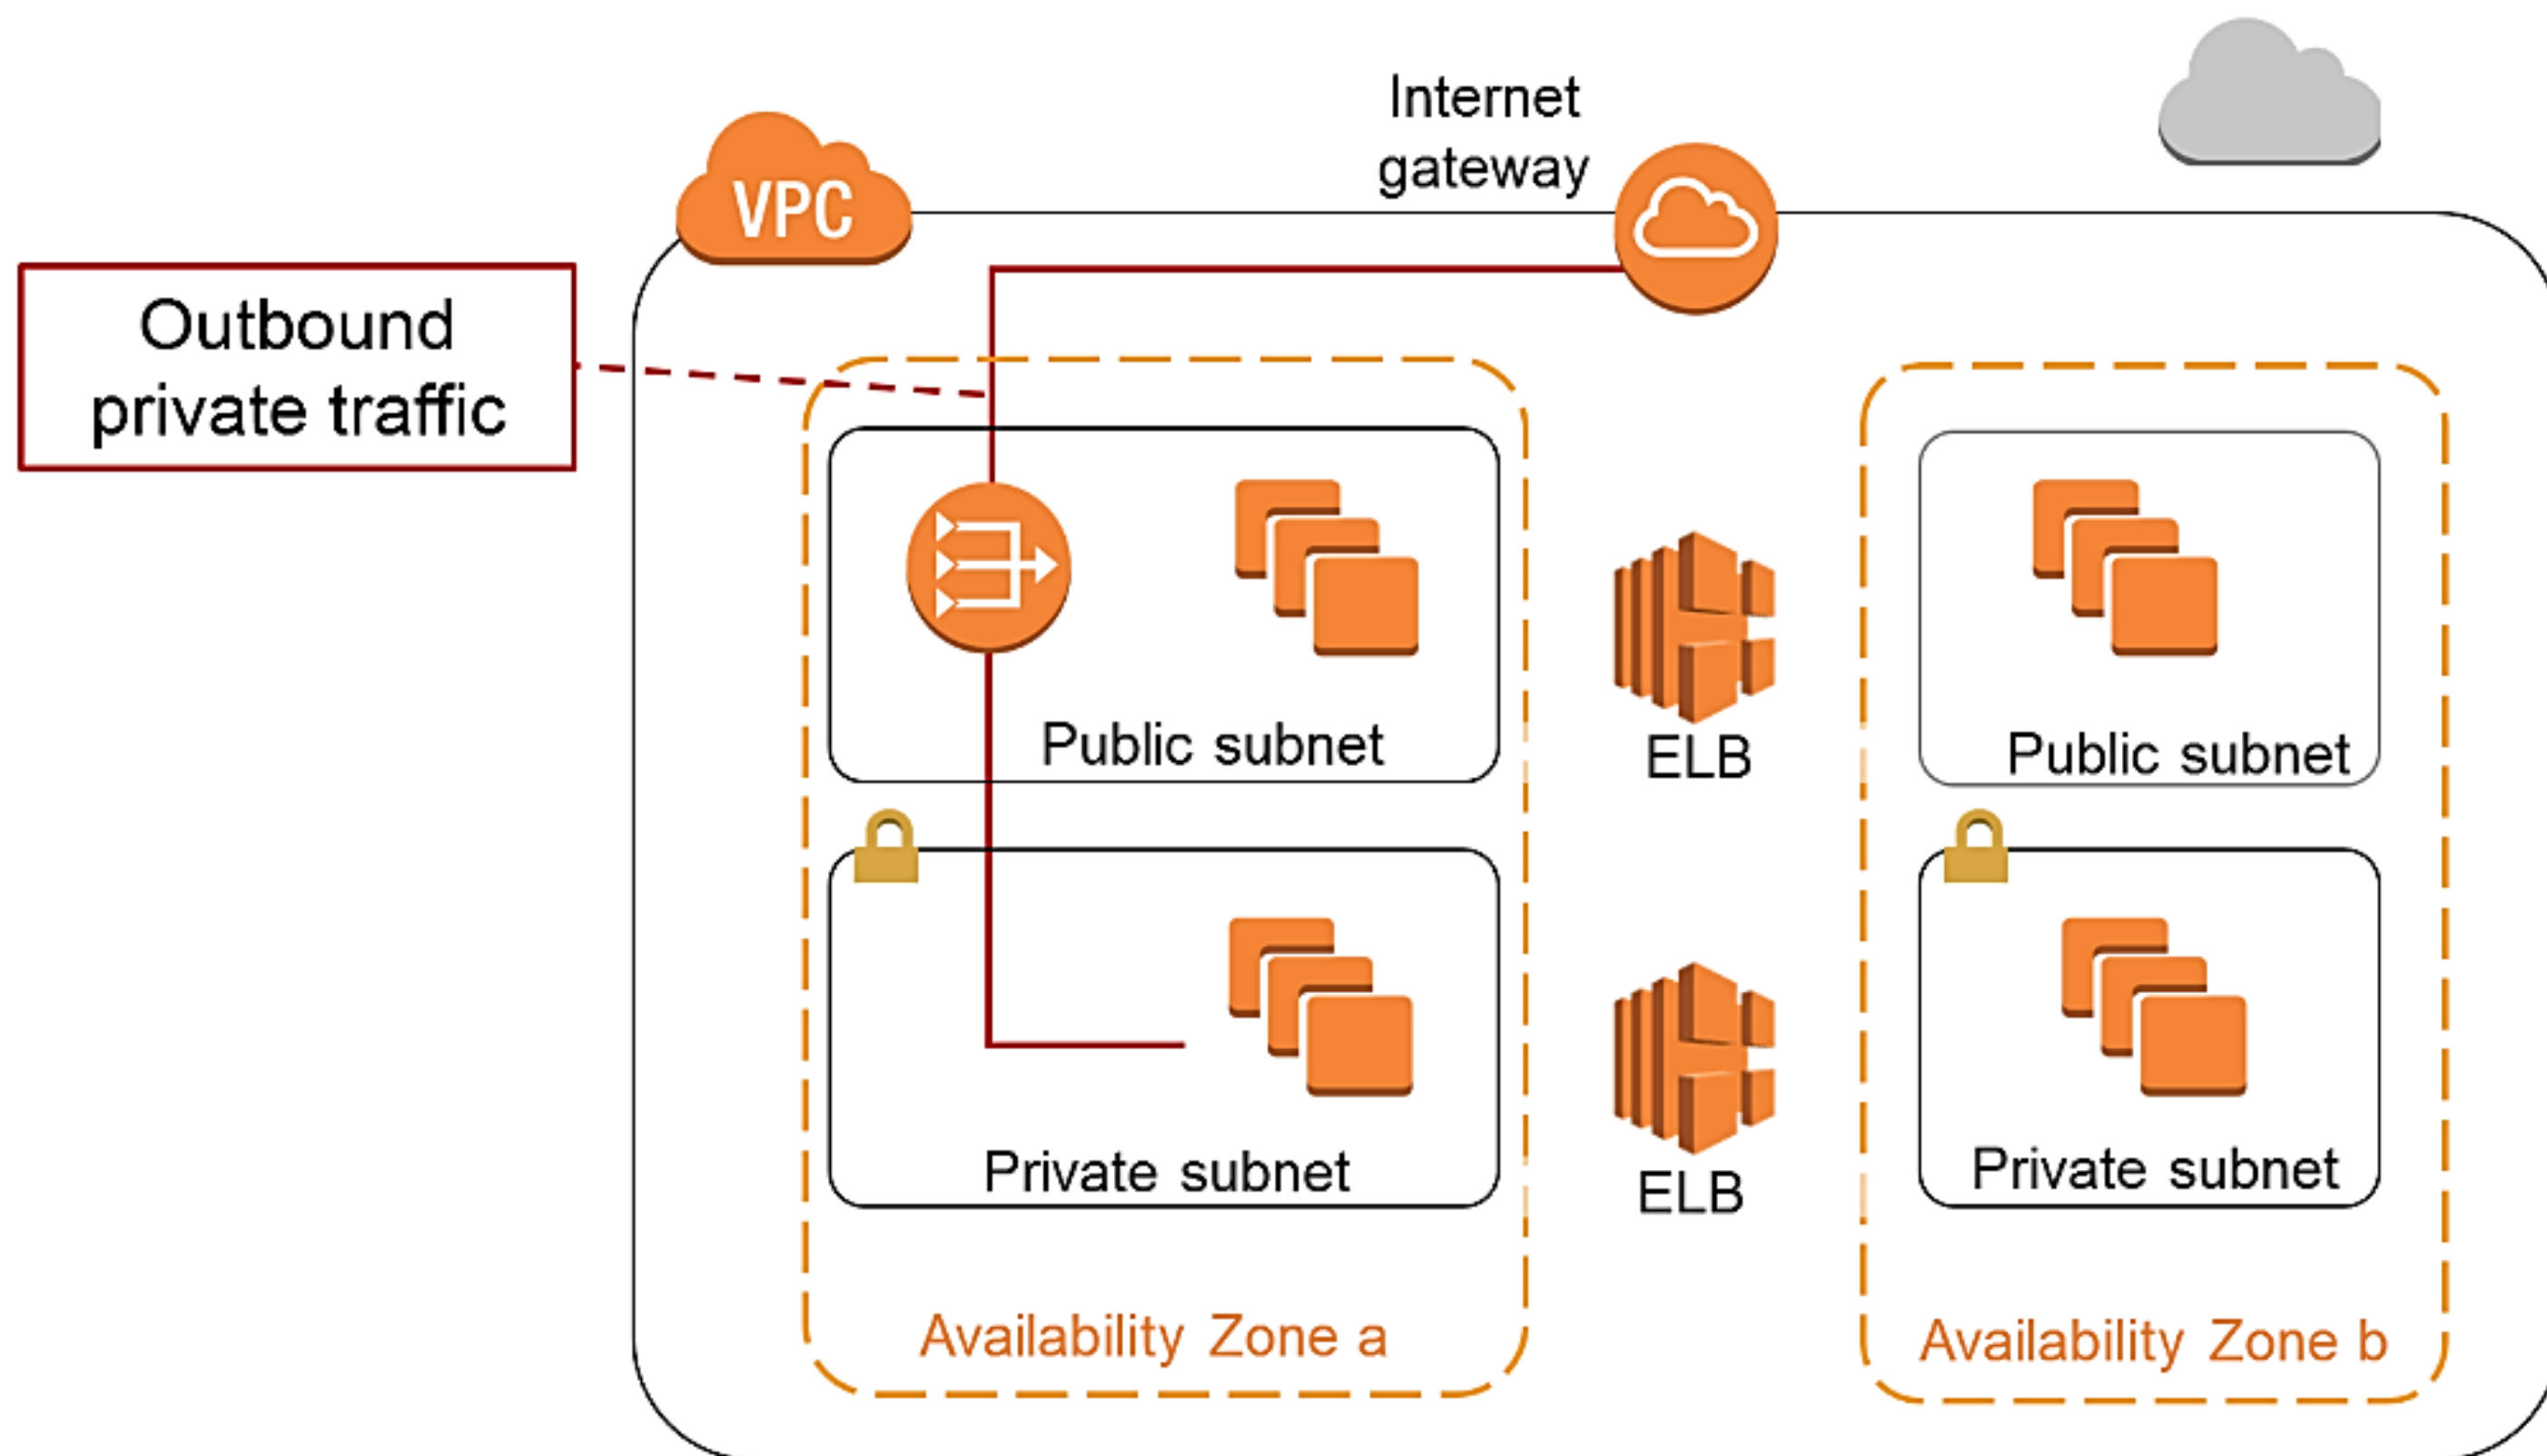
Start with two Availability Zones per AWS Region.

If resources in one Availability Zone are unreachable, your application shouldn't fail.













Auto Scaling



- Launches or terminates instances based on specified conditions
- Automatically registers new instances with load balancers when specified
- Can launch across Availability Zones
- Additional AutoScaling instead of EC2: Application AS (ECS, EMR, DynamoDB), AWS AutoScaling (Managed Both).






- Desired capacity
- Minimum capacity
- Maximum capacity

What would be a good **minimum** capacity to set it to?

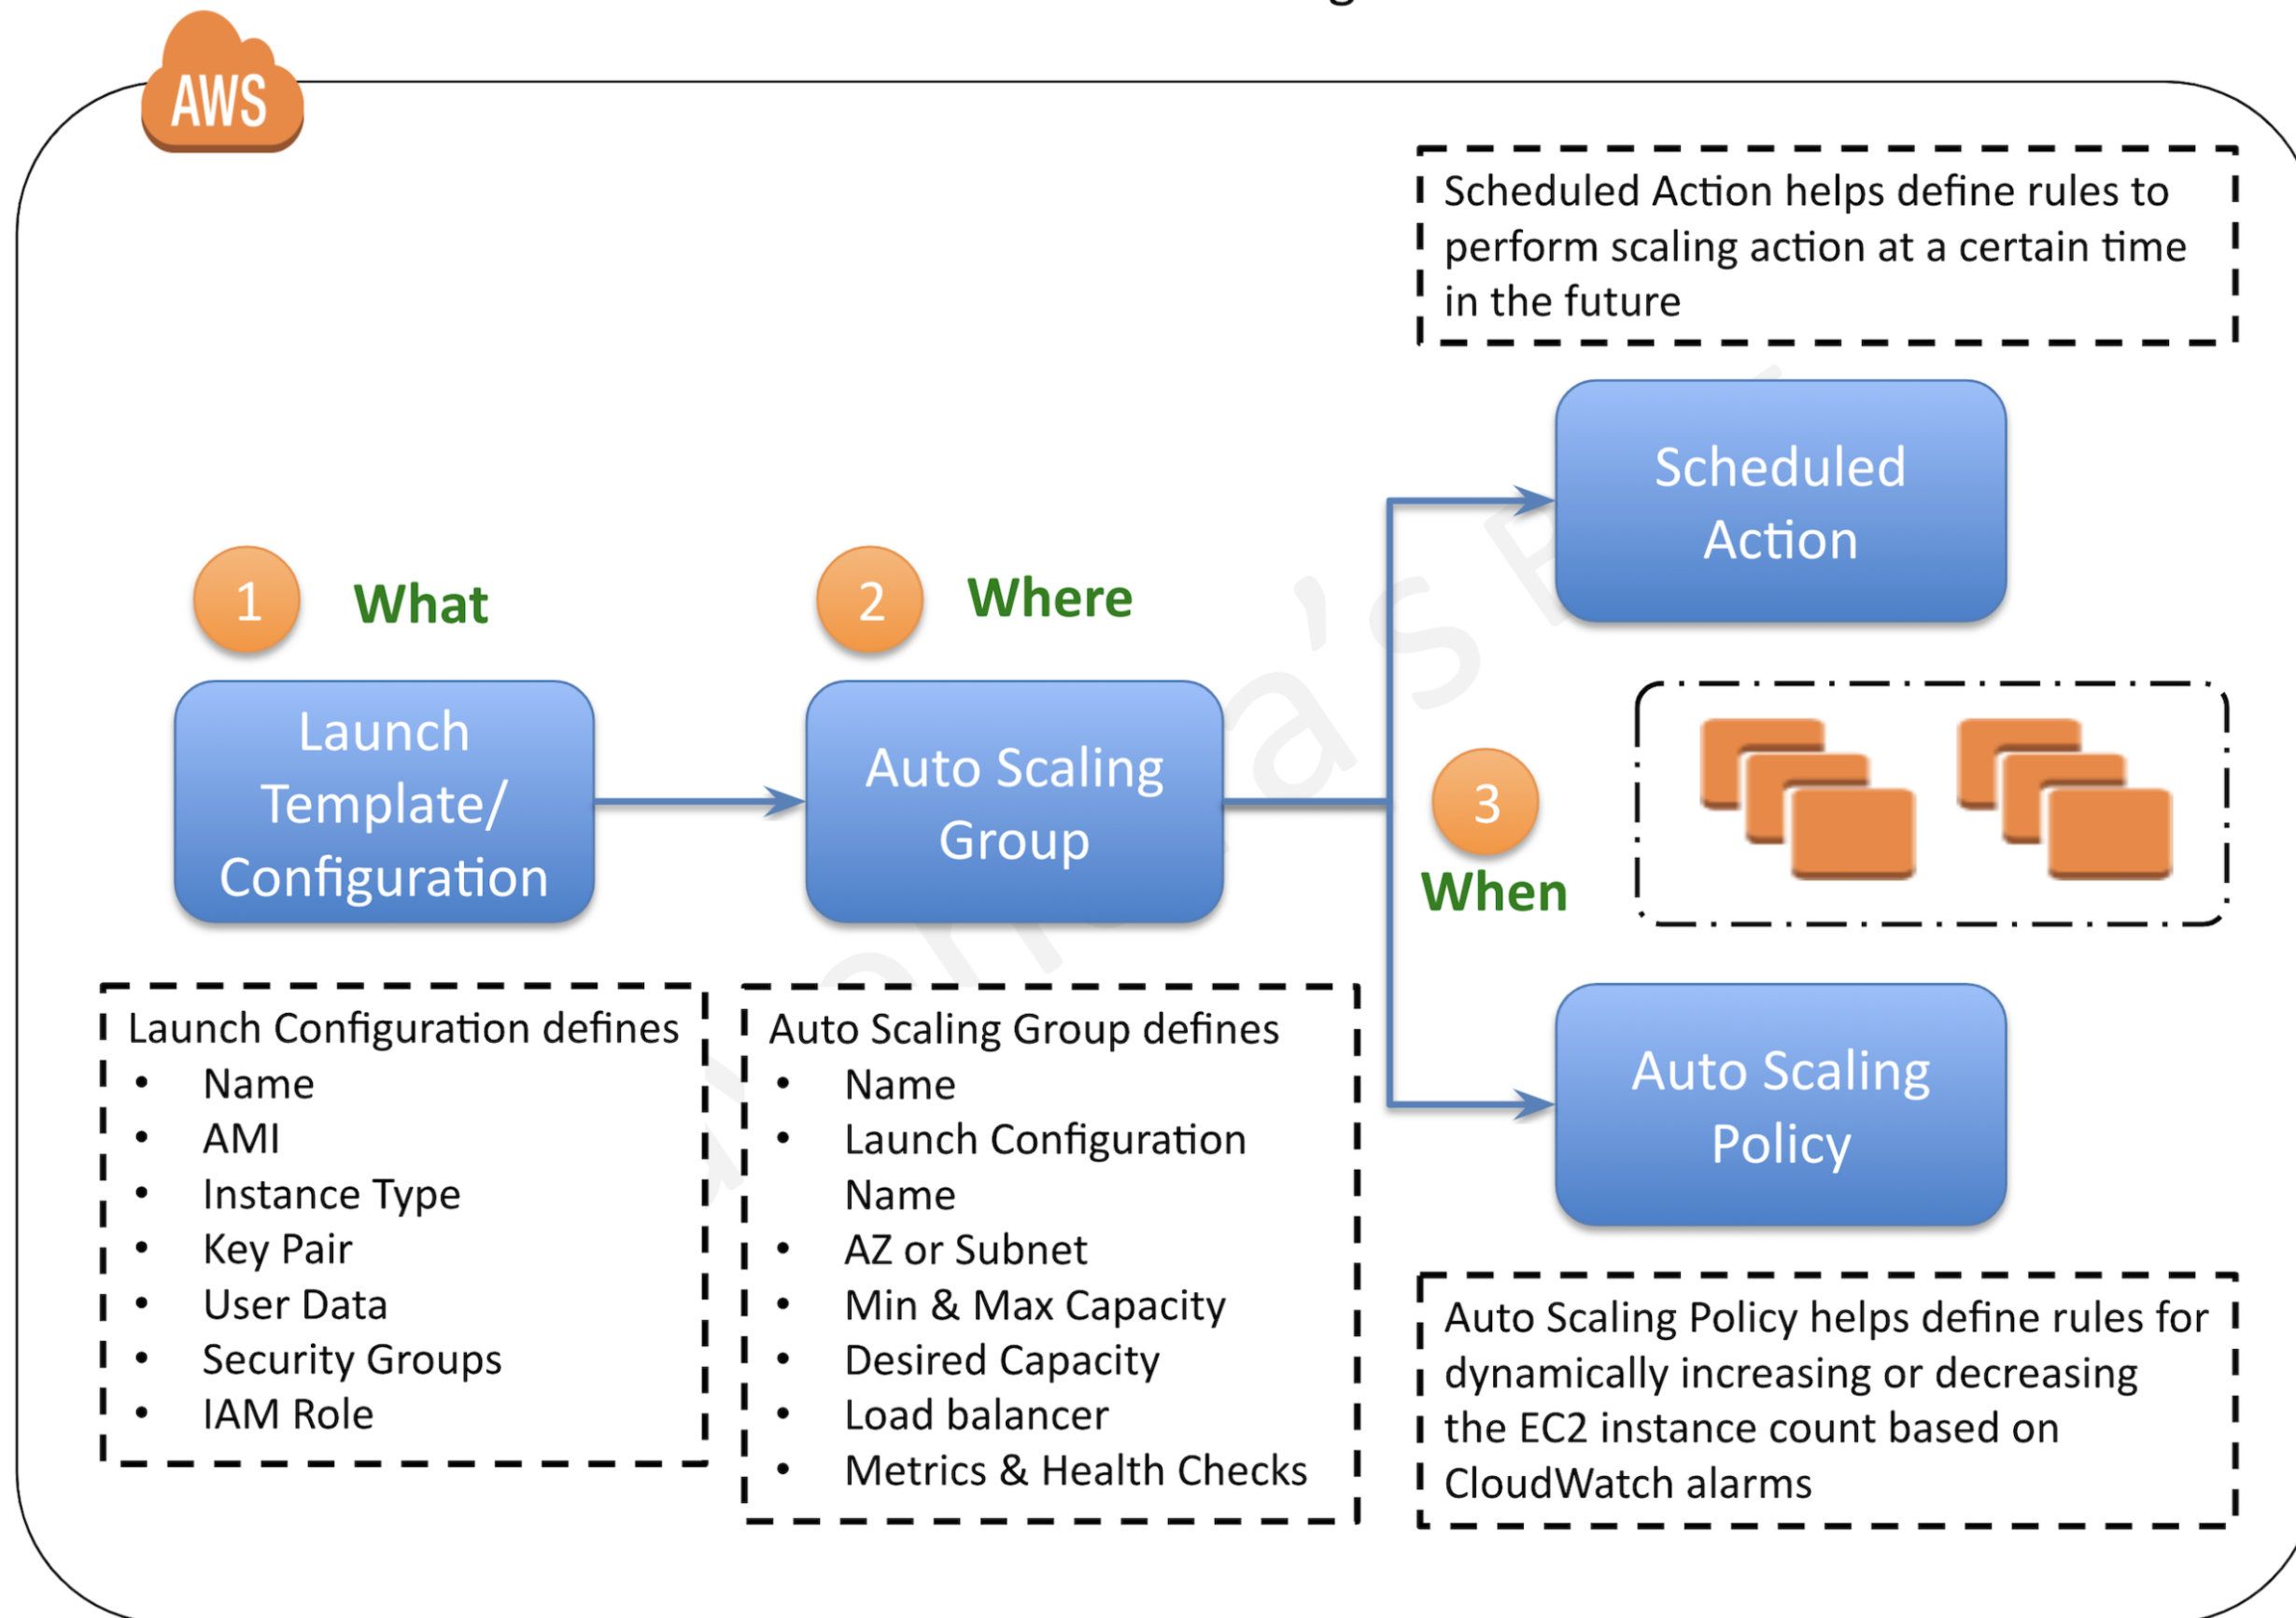
What would be a good **maximum** number?

By default, **Desired** (# instances all the time) is equal to **minimum**.





AWS Auto Scaling

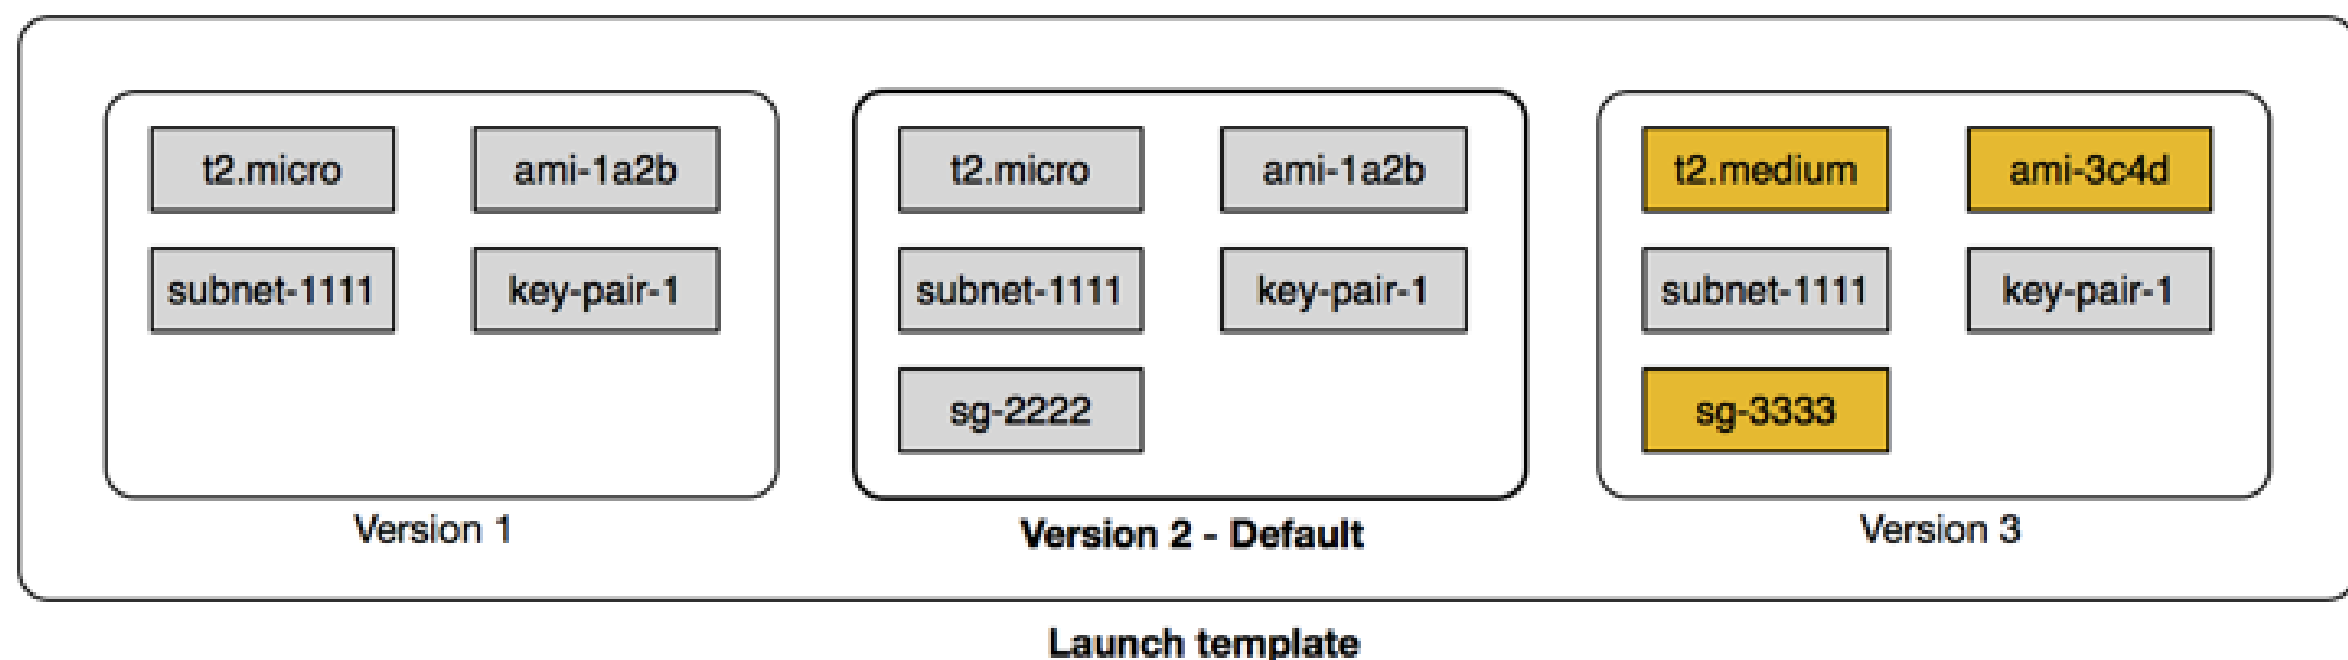




Launch template is a better option than Launch Configuration because it allows **reused** and add configuration to new templates.

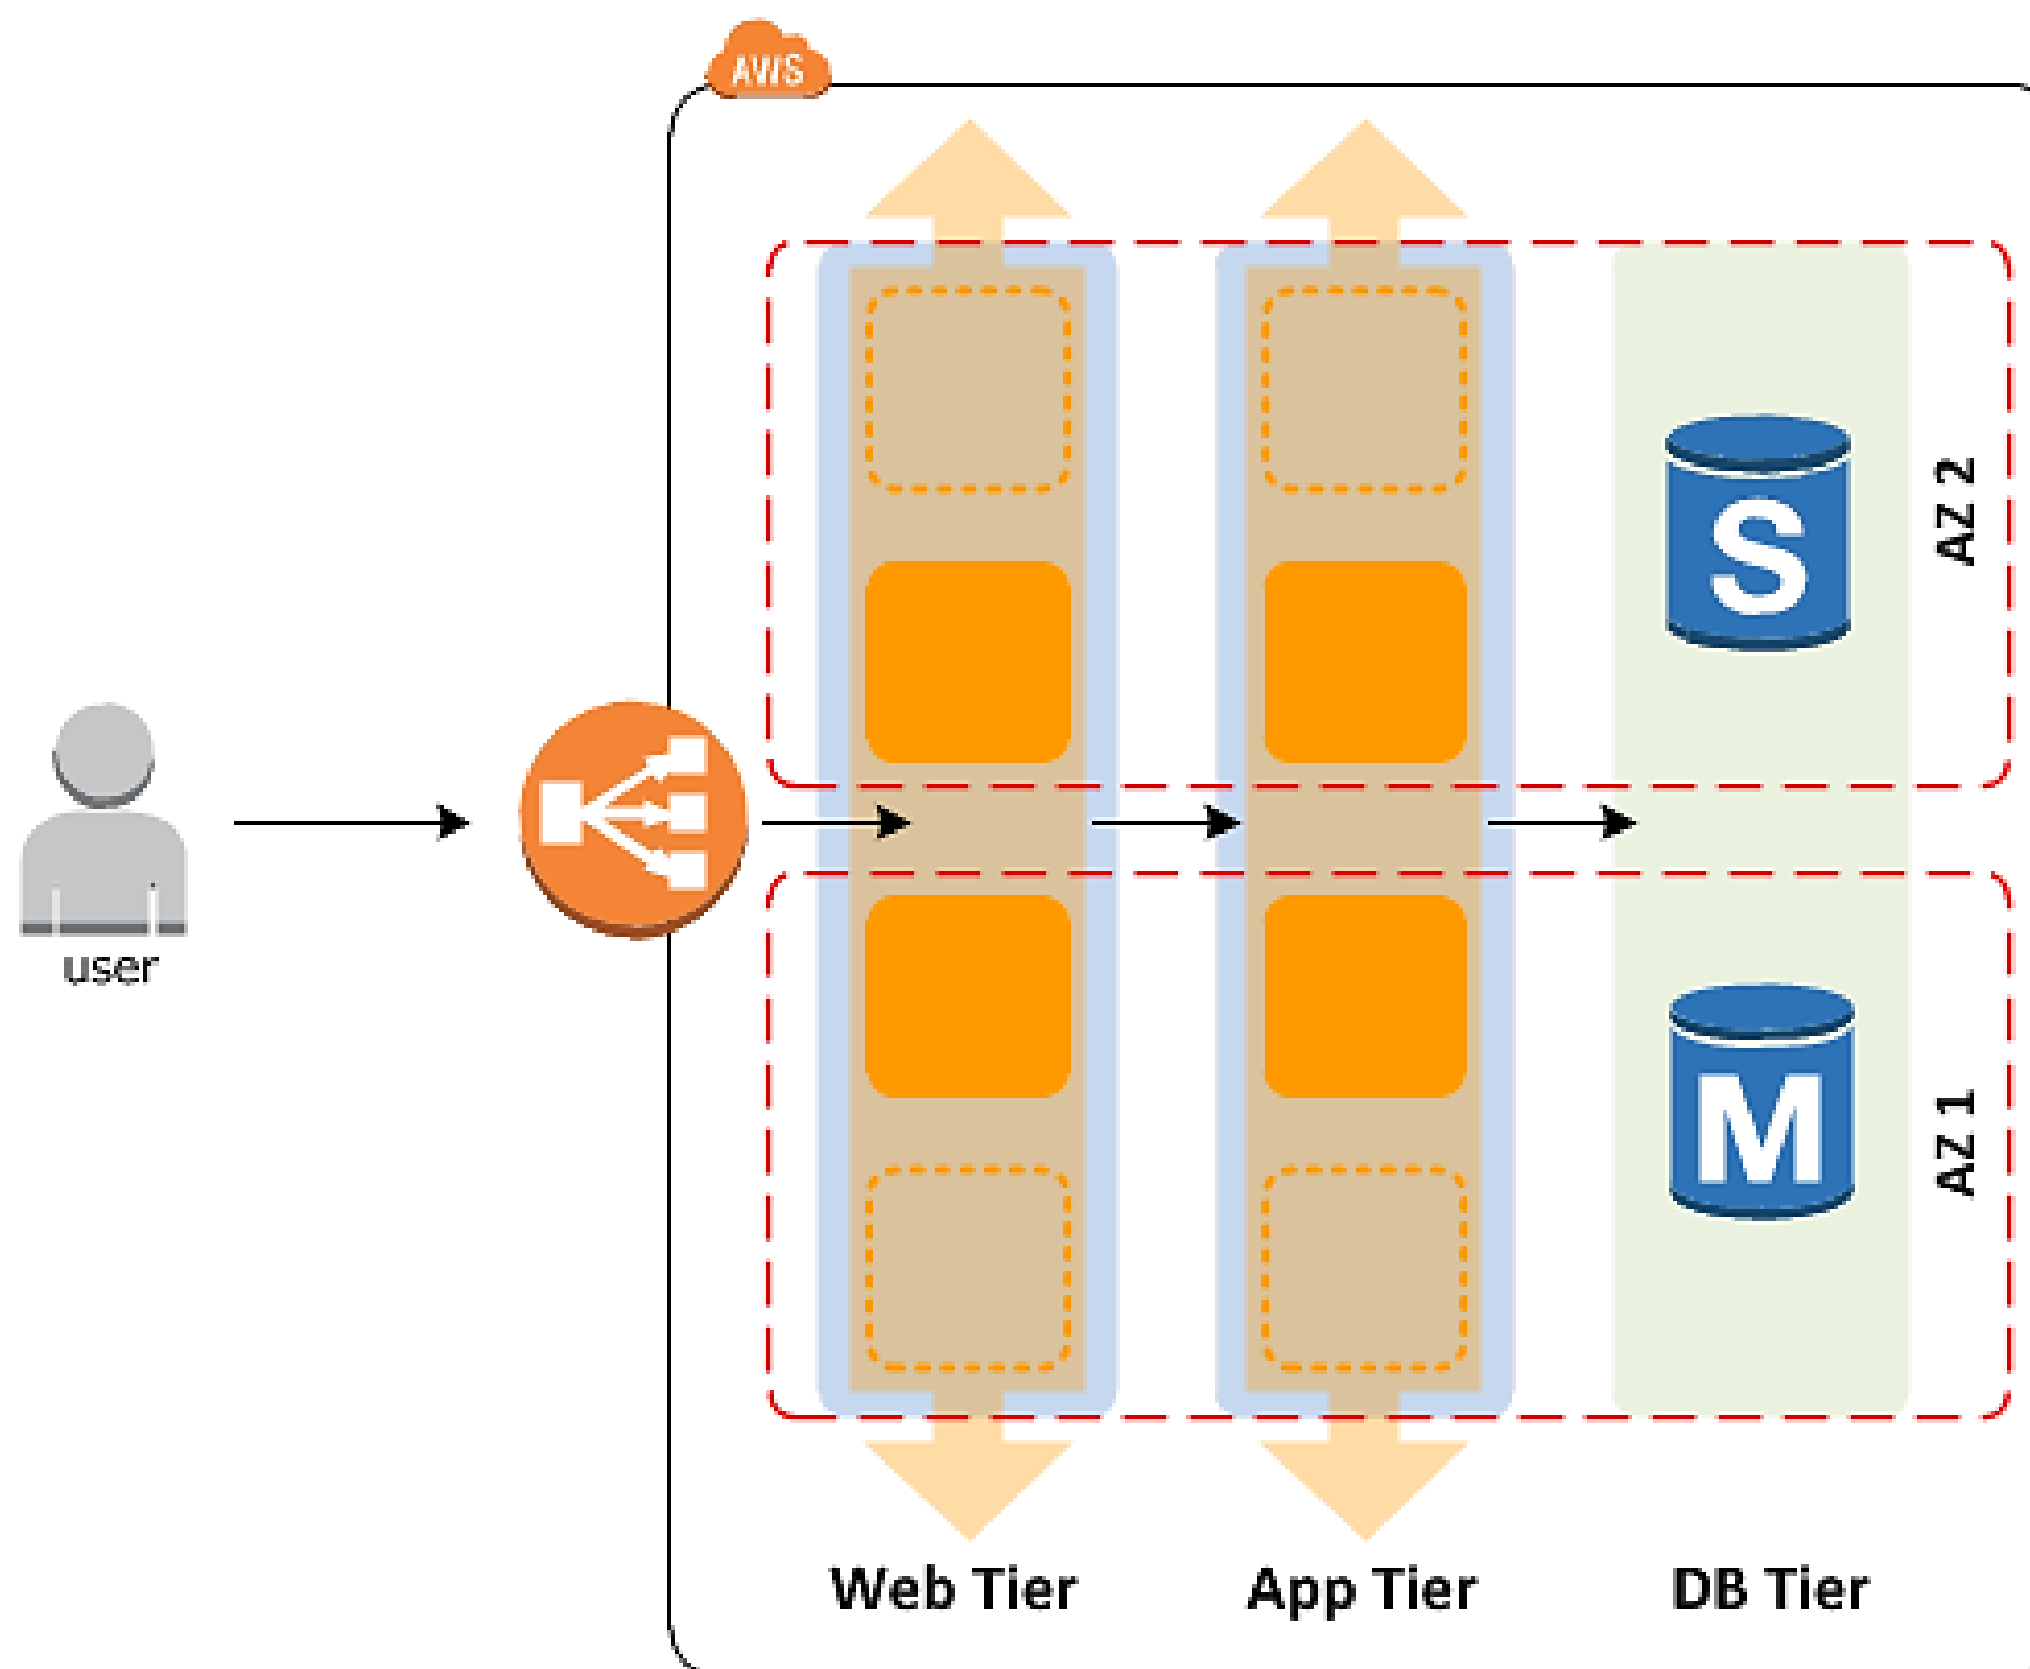
In addition, launch templates allows all configuration from AutoScaling groups, for instance mix OnDemand and Spot Instance on the same group or launch on Dedicated Hosts.


Both options contains the same information about instances: type, AMI, Keypair, Security Group, storage, etc.





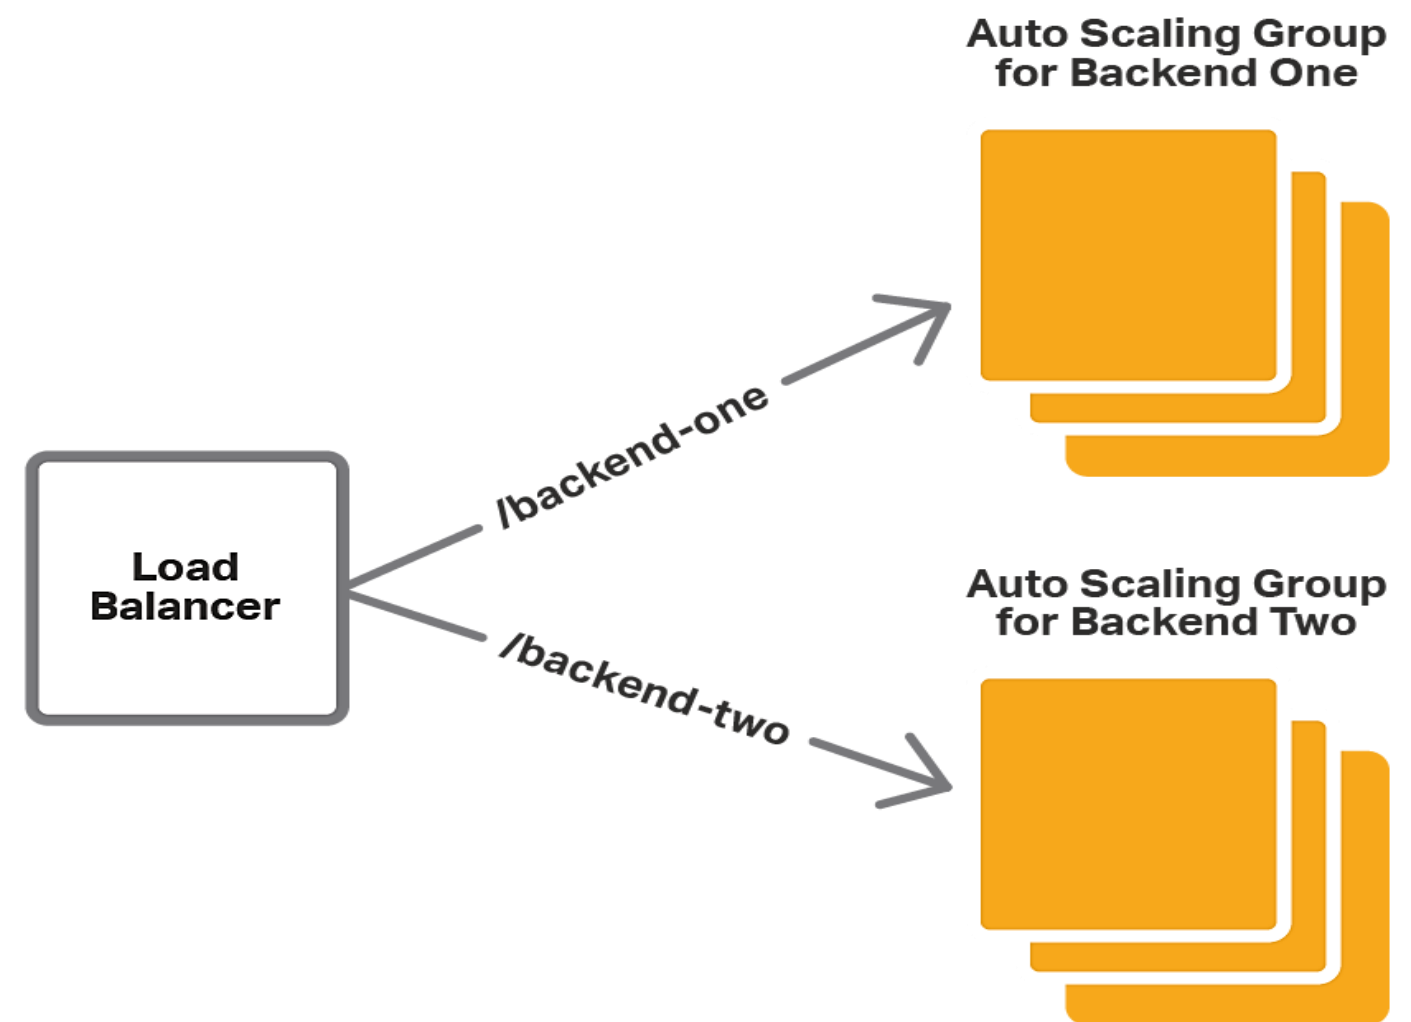
ASG Components





Policies:

- Manual
- Scheduled
- Dynamic
- Dynamic: Target Tracking Policy
- Dynamic: Simple Tracking Policy → Wait until health check and cooldown before evaluate
- Dynamic: Step Tracking Policy
- Metrics on CloudWatch apply on Dynamic.





1. Choose LB (i.e. SSL Termination, Global Balancing with R53)
2. Choose the appropriate instance families and sizing based on the workload in that ASG
3. Consider placement groups in ASG –if applicable-
4. Use Launch Templates
5. Group instances by purpose (i.e. Domains)
6. Set up health checks
7. Utilize target tracking scaling policies
8. Implement cooldown periods
9. Configure notifications
10. Implement proper security (NACLs? Sec Groups)
11. Implement instance termination policies
12. Distribute instances across Availability Zones
13. Use Auto Scaling lifecycle hooks (Launch/Terminate)



Amazon CloudWatch



Alarm



Rule



Event
(event-based)



Event
(time-based)

- Monitoring and Observability Service to cover AWS Resources, especially computing services such Applications and Containers, in addition Logs, Databases. It collects and track metrics from those resources.
- You can create: Dashboards, Insights, Alarms and Events. You collect: Metrics and Logs.
- Services associated: Cloud Trail, IAM
- Services to trigger actions: SNS – Simple Notification Service, AutoScaling Groups, Lambda; however for Events you can trigger more services.



Metrics:

Metrics: Timestamp (UTC), Retention

High-Resolution Monitoring: - Retention: 3 h.

Detailed monitoring: 1 mins (Paid) – R: 15 d.

Basic monitoring: 5 mins – R: 63 d.

Statistics: Max, Min, Average, Sum, Count, Percentiles

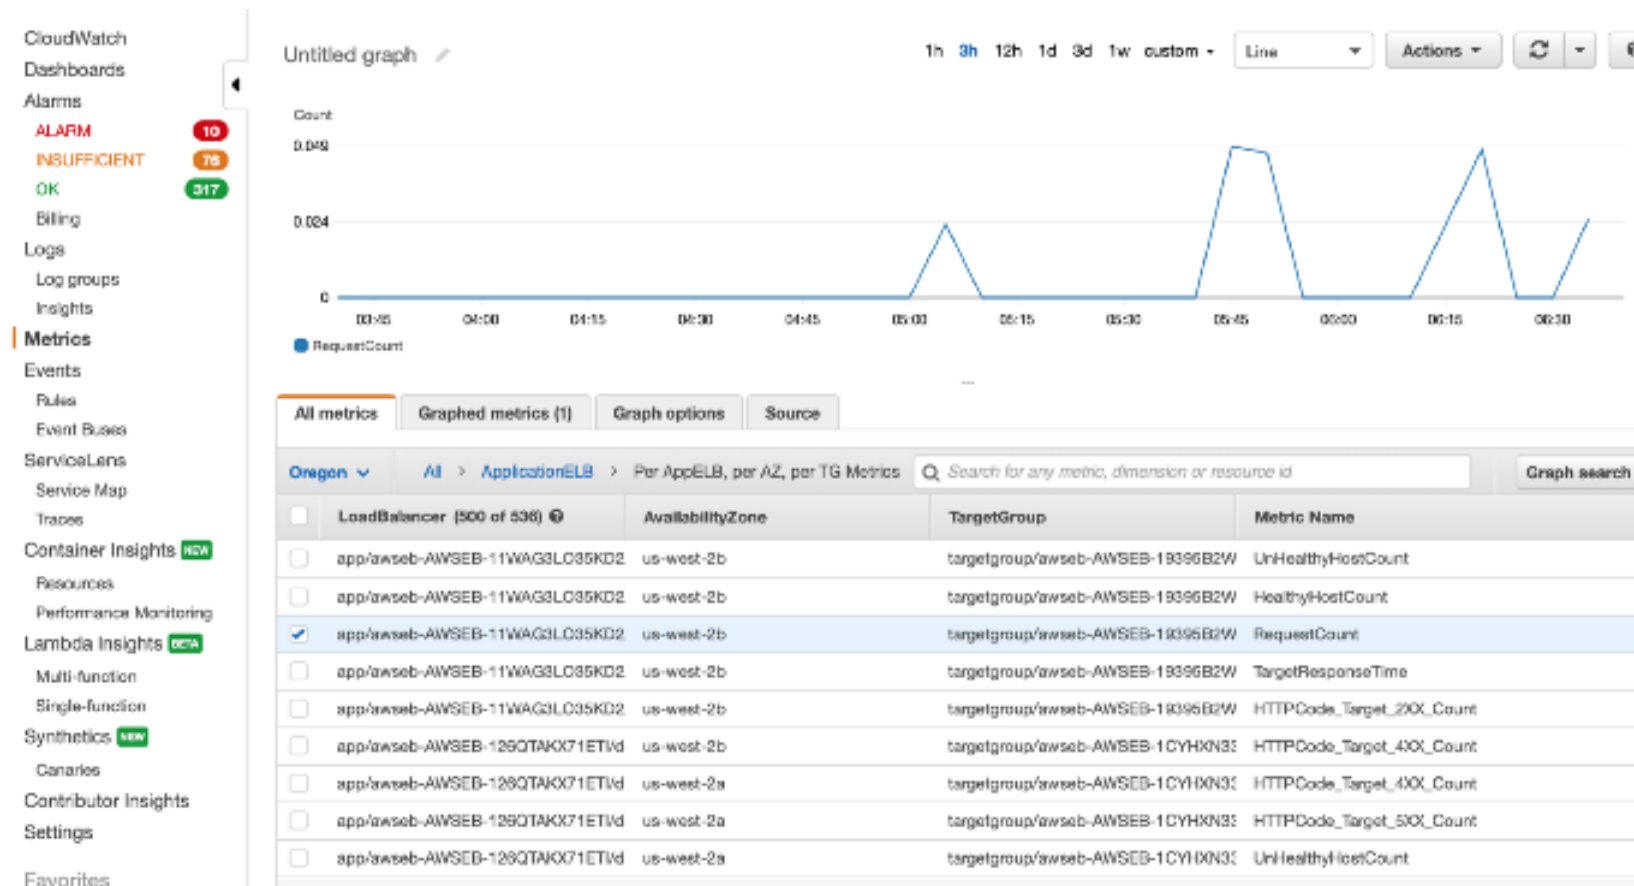
Units, Period

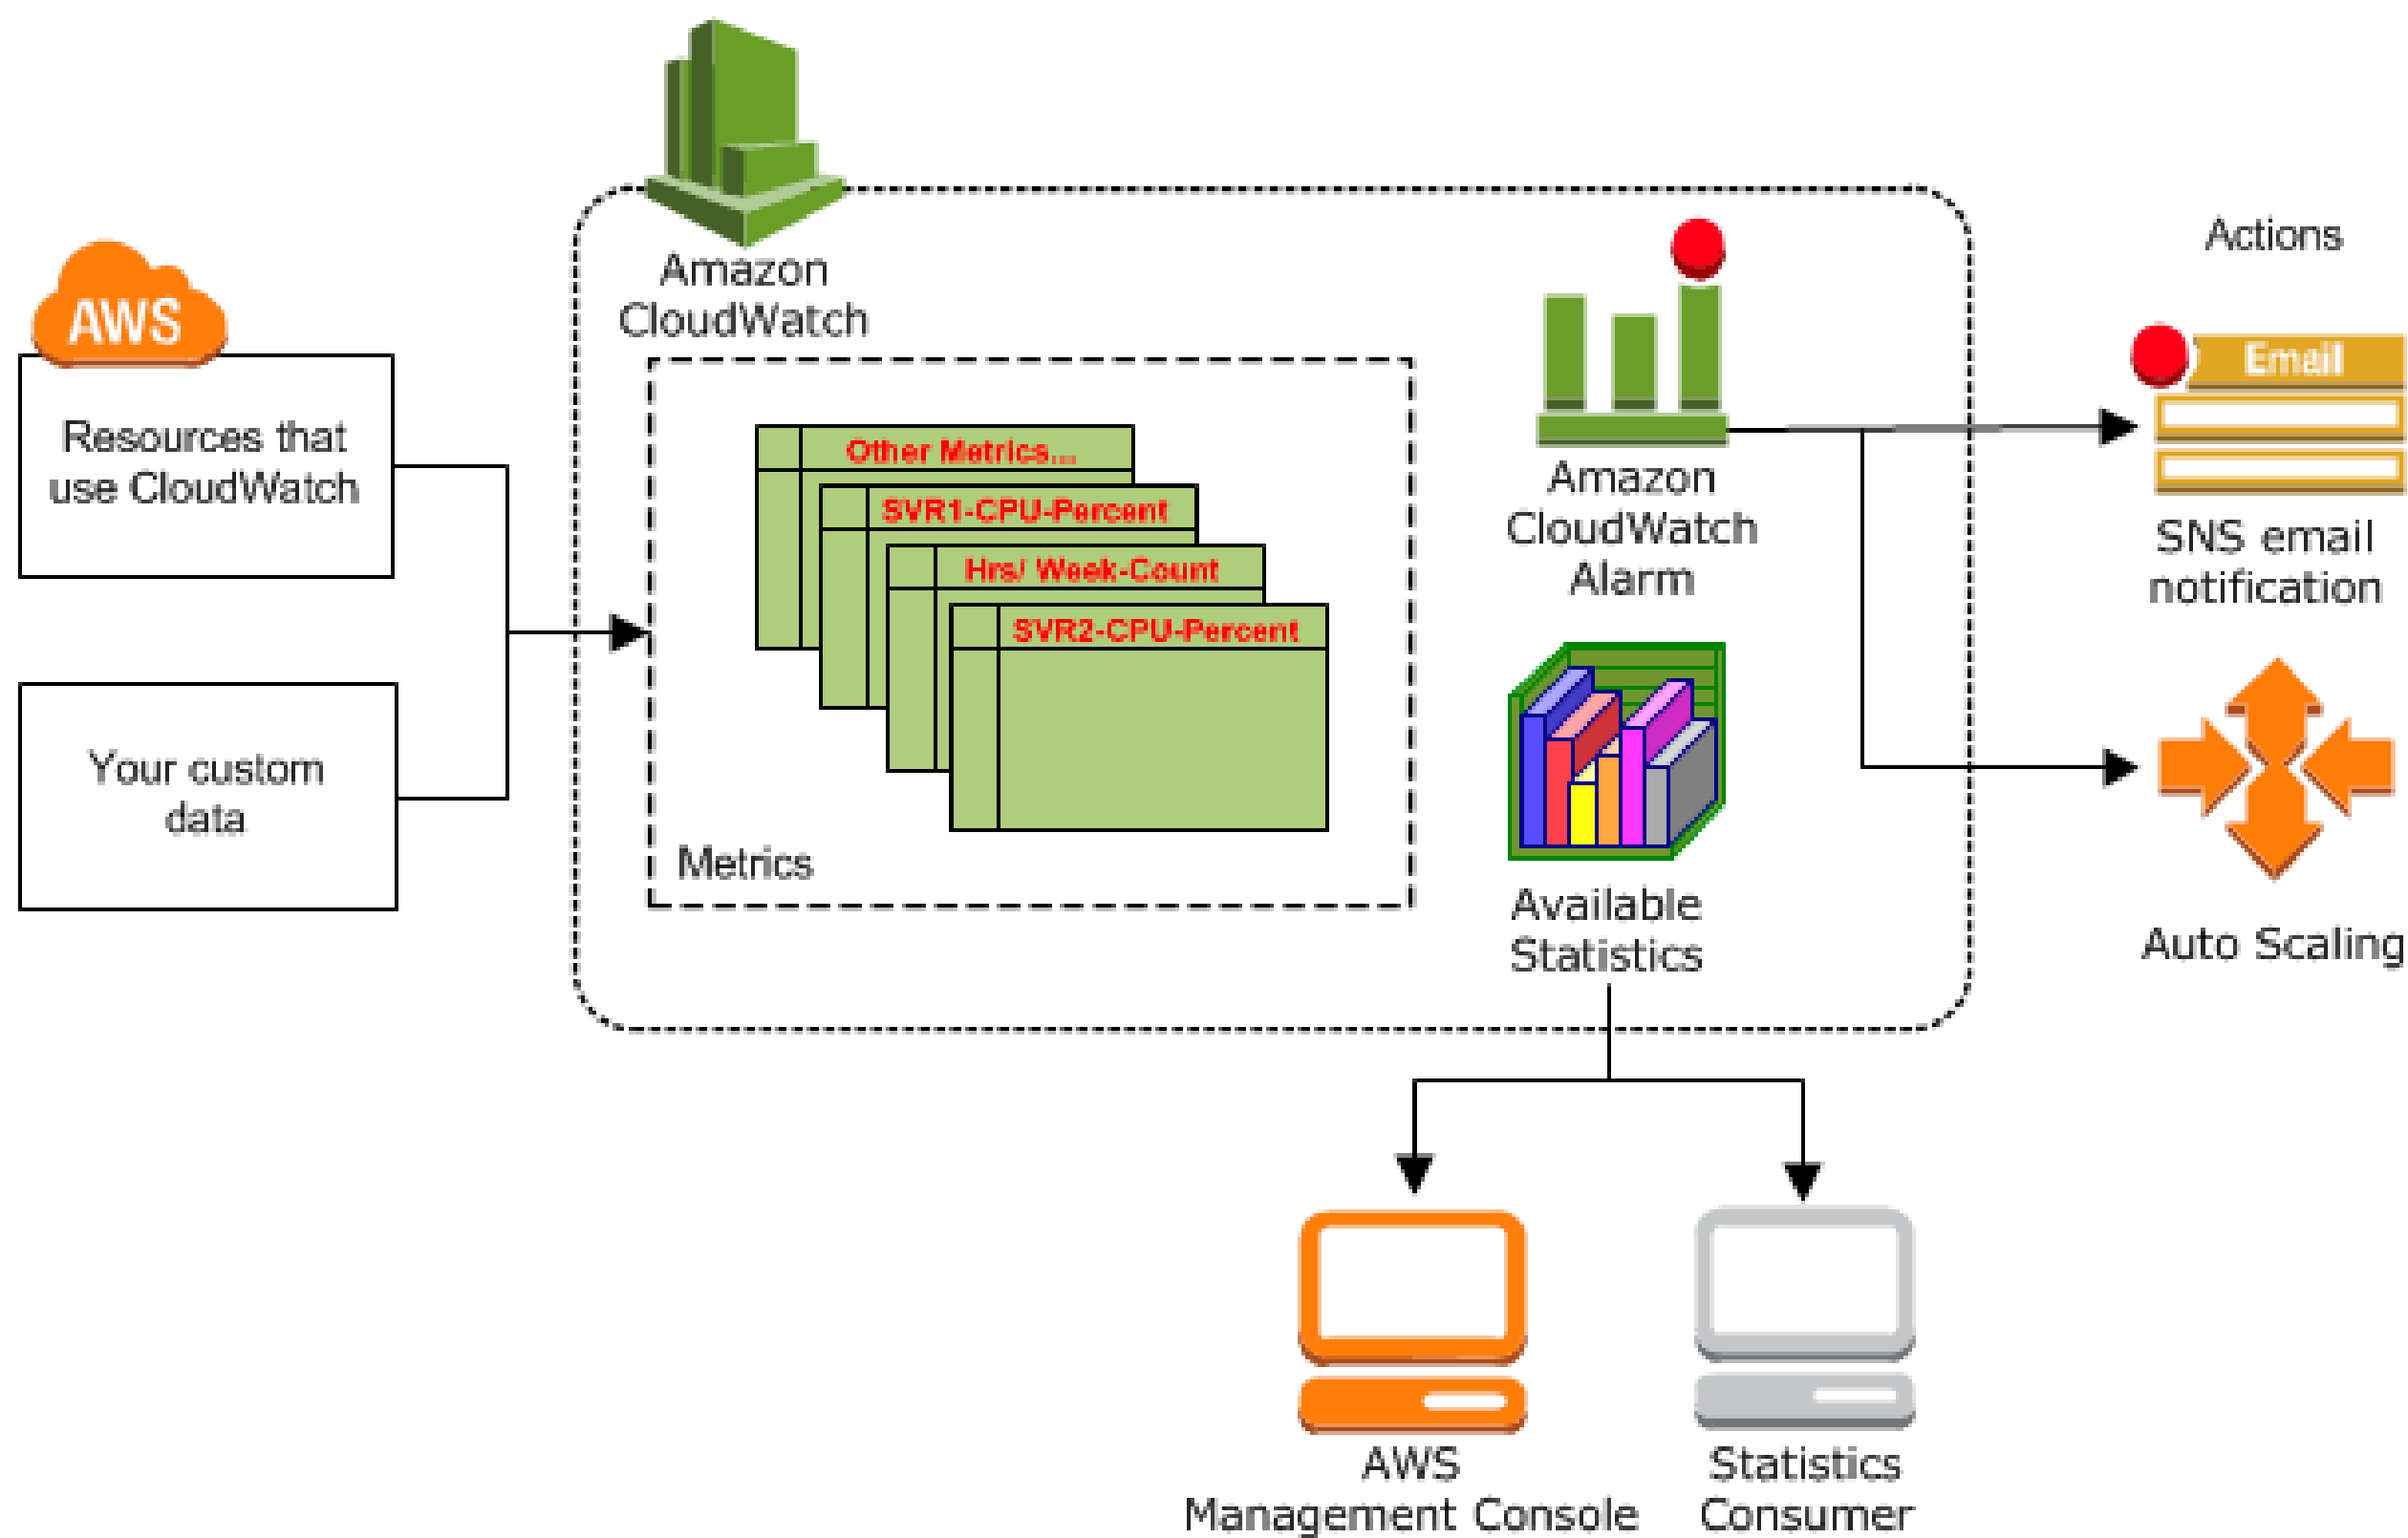
Aggregation

Concepts: Namespaces (i.e. Service) / Dimensions (i.e. Group Metrics)

Alarms: Trigger SNS and/or AutoScaling Groups

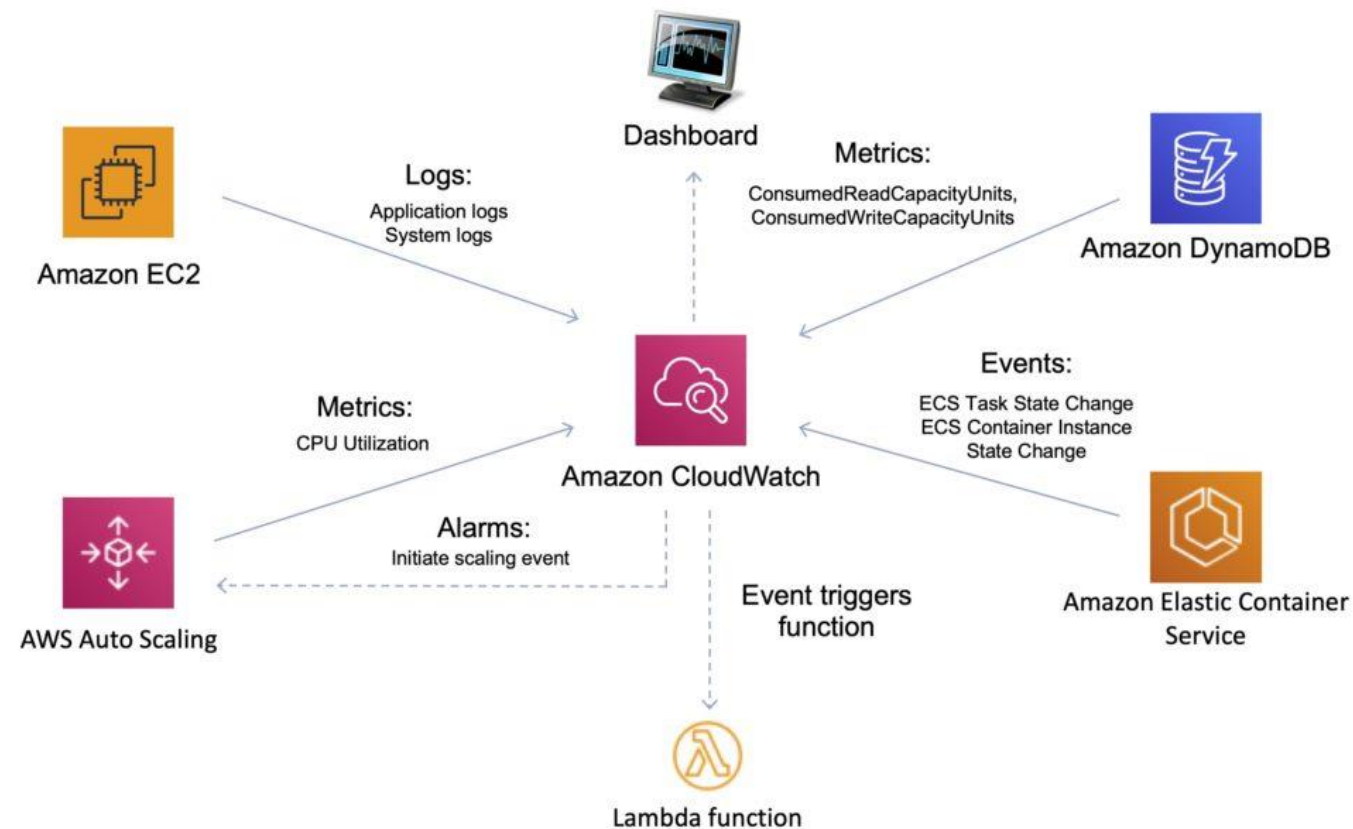
Events: Trigger something by Event or Schedule








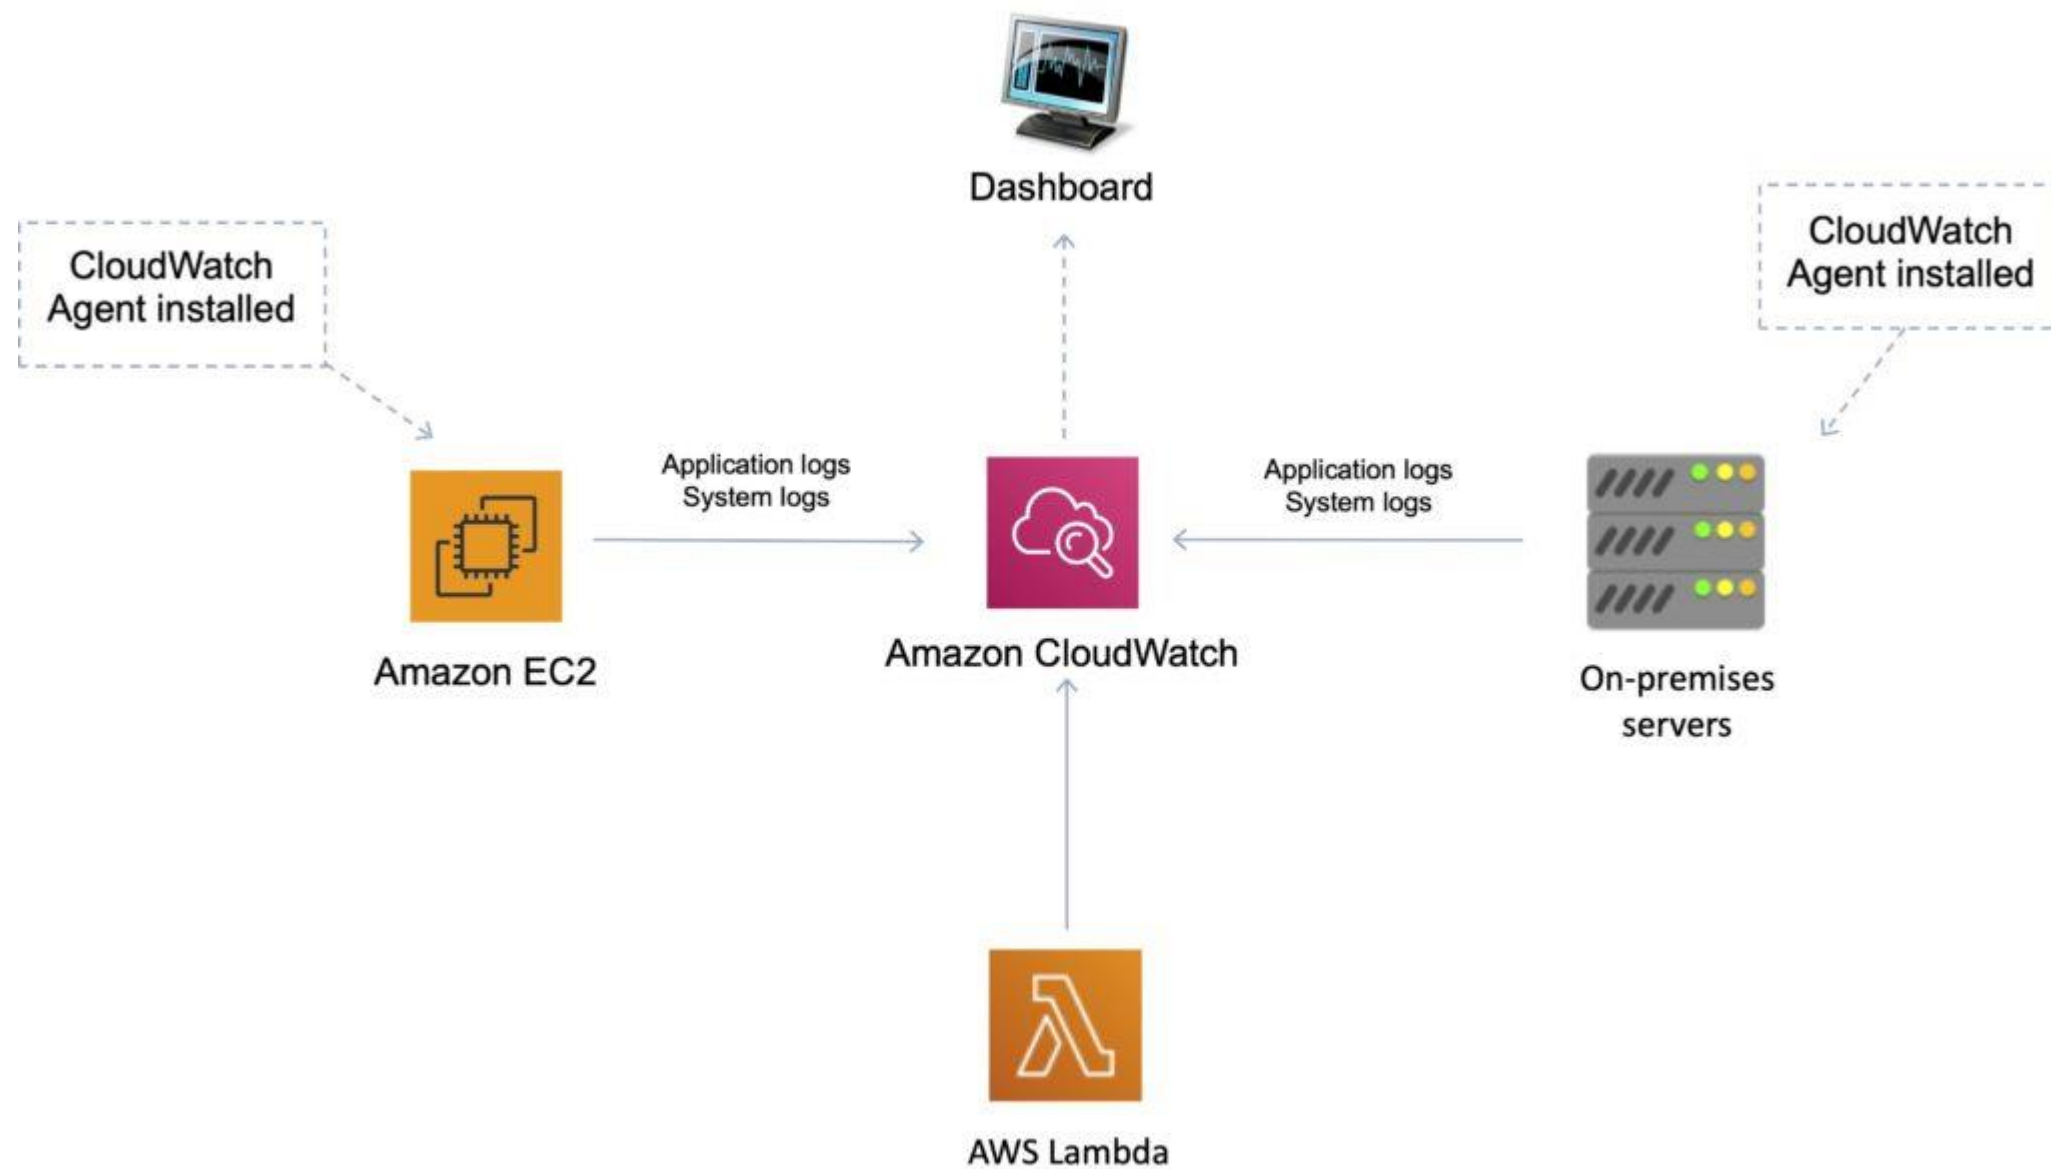
CloudWatch Alarms
use conditions to trigger
notifications/actions



CloudWatch Events
use rules to trigger targets



Alarms	Events
Can be trigger by events/sustained events	Can be trigger by events or timing
Limited number of actions: SNS, AutoScaling Group, EC2 Actions.	Many services to trigger, which called Targets: Lambda, SNS, etc.
Watch Single Metric (Defined by a period)	Respond to Actions (Near real-time)
Can be added to Dashboard	No can be added to Dashboard



It receives Application and System Logs, and act as centralized repository and then, it can be feed by a Lambda or ElasticSearch Cluster.